

# Chapter 7

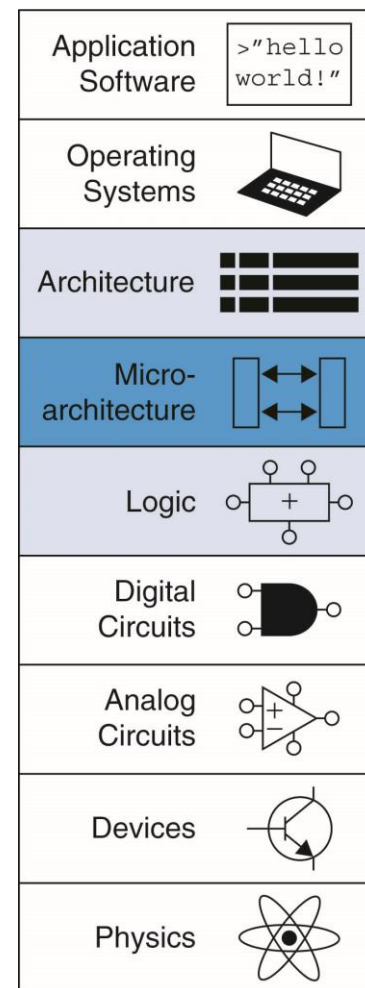
## ***Digital Design and Computer Architecture, 2<sup>nd</sup> Edition***

---

David Money Harris and Sarah L. Harris

# Chapter 7 :: Topics

- Introduction
- Performance Analysis
- Single-Cycle Processor
- Pipelined Processor
- Exceptions
- Advanced Microarchitecture



# Introduction

- **Microarchitecture:** the implementation of an architecture in hardware
- **Processor:**
  - **Datapath:** functional blocks
  - **Control:** control signals

Application Software	programs
Operating Systems	device drivers
Architecture	instructions registers
Micro-architecture	datapaths controllers
Logic	adders memories
Digital Circuits	AND gates NOT gates
Analog Circuits	amplifiers filters
Devices	transistors diodes
Physics	electrons



# Microarchitecture

- Multiple implementations for a single architecture:
  - **Single-cycle:** Each instruction executes in a single cycle
  - **Multicycle:** Instructions are broken into series of shorter steps. Each instruction executes in  $n$  cycles, where  $n$  varies according to the instr.
  - **Pipelined:** Each instruction broken up into series of steps & multiple instructions execute at once (Note: AMD and Intel pipelines are different, for the same IA-32 architecture (a.k.a. x86 ISA))



# Processor Performance

- Program execution time

Execution Time = (#instructions)(cycles/instruction)(seconds/cycle)

- Definitions:
  - IC: Instruction Count (= #instructions)
  - CPI: Cycles/Instruction
  - clock period: seconds/cycle
  - IPC: Instructions/Cycle (= 1/CPI)
- Challenge is to satisfy constraints of:
  - Cost
  - Power
  - Performance

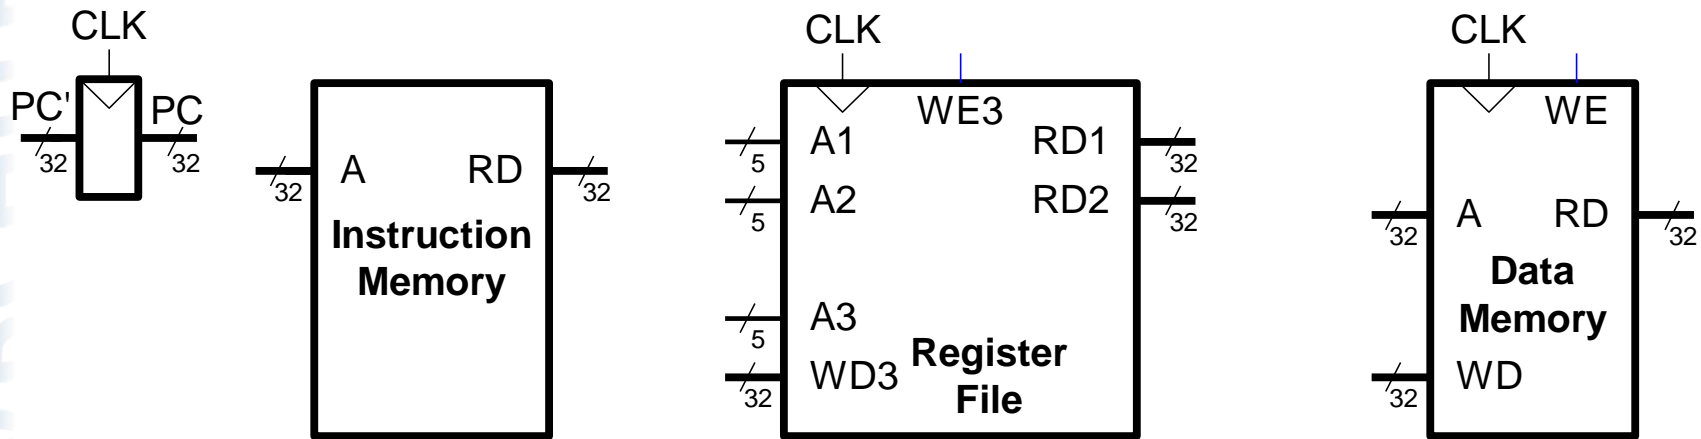
# MIPS Processor

- Consider subset of MIPS instructions:
  - R-type instructions: `and`, `or`, `add`, `sub`, `sllt`
  - Memory instructions: `lw`, `sw`
  - Branch instructions: `beq`

# Architectural State

- Determines everything about a processor:
  - PC and special registers
  - Register File
  - Memory

# MIPS State Elements



Plus the HI and LO registers

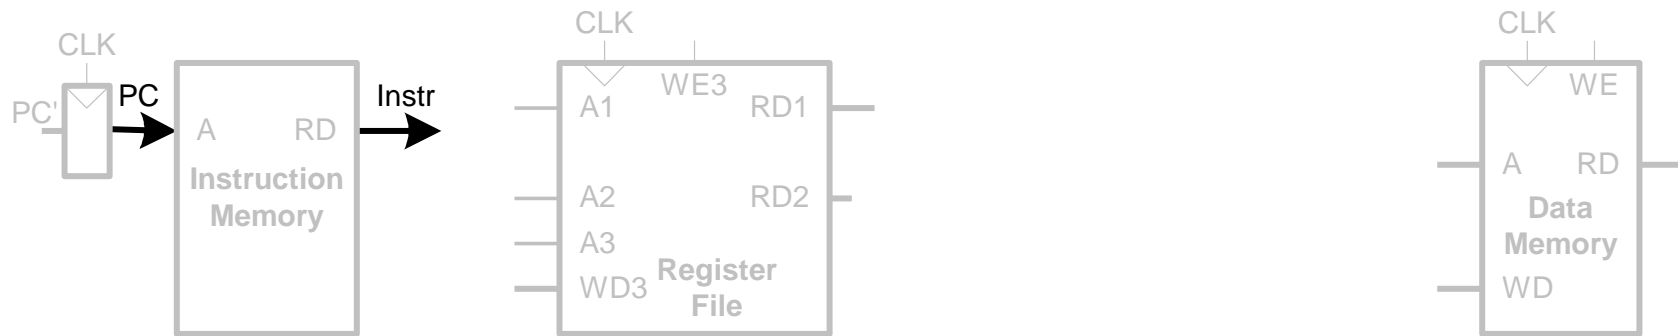


# Single-Cycle MIPS Processor

- Datapath—design it 1<sup>st</sup>, to make the instruction actions possible
- Control—design it 2<sup>nd</sup>, to make them happen


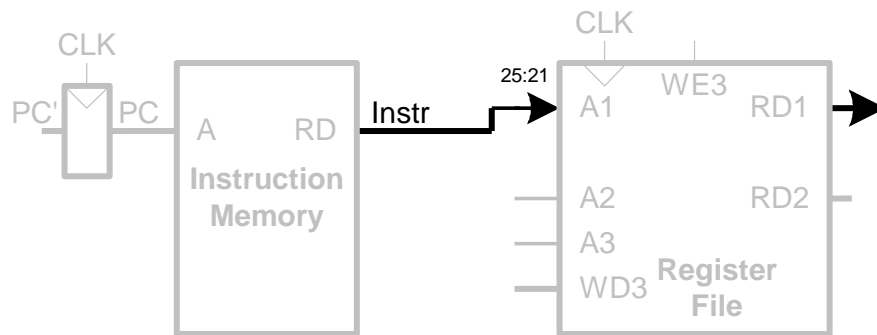
# Single-Cycle Datapath: $1_w$ fetch

## STEP 1: Fetch instruction $IM[PC]$



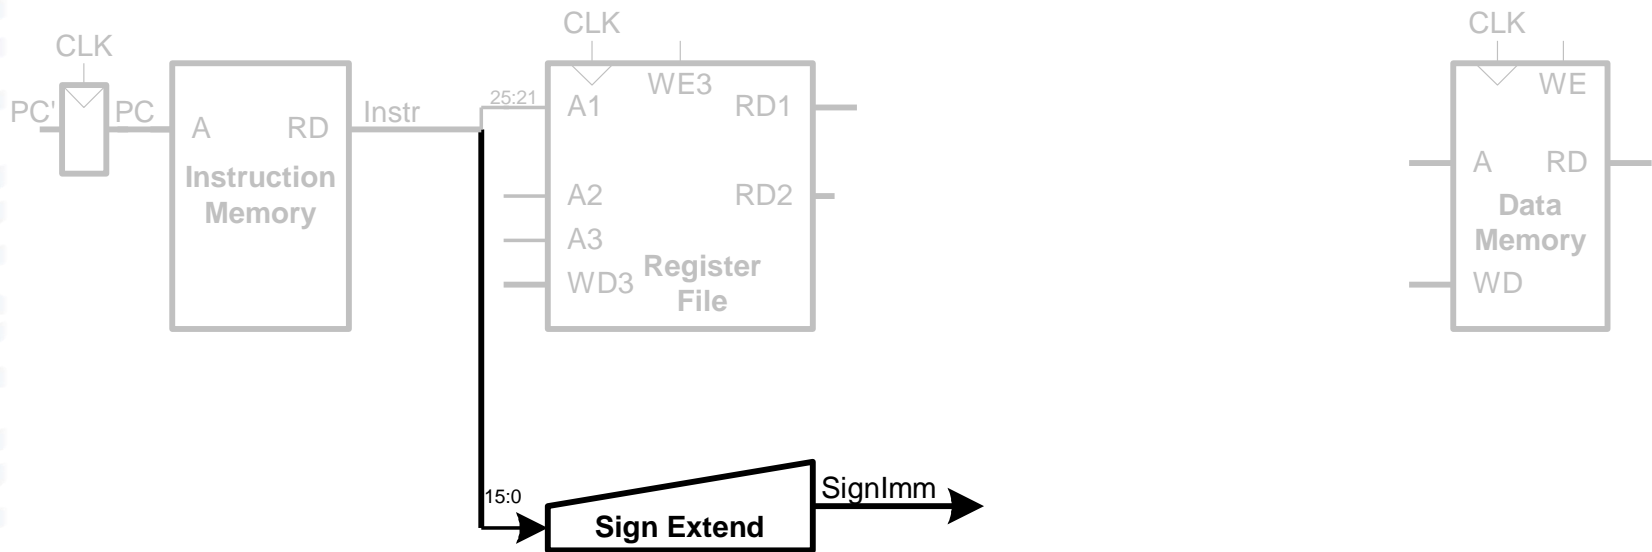
# Single-Cycle Datapath: $1_w$ Register Read

**STEP 2:** Read source operands from RF  
RF[rs] or RF[Instr(25:21)]



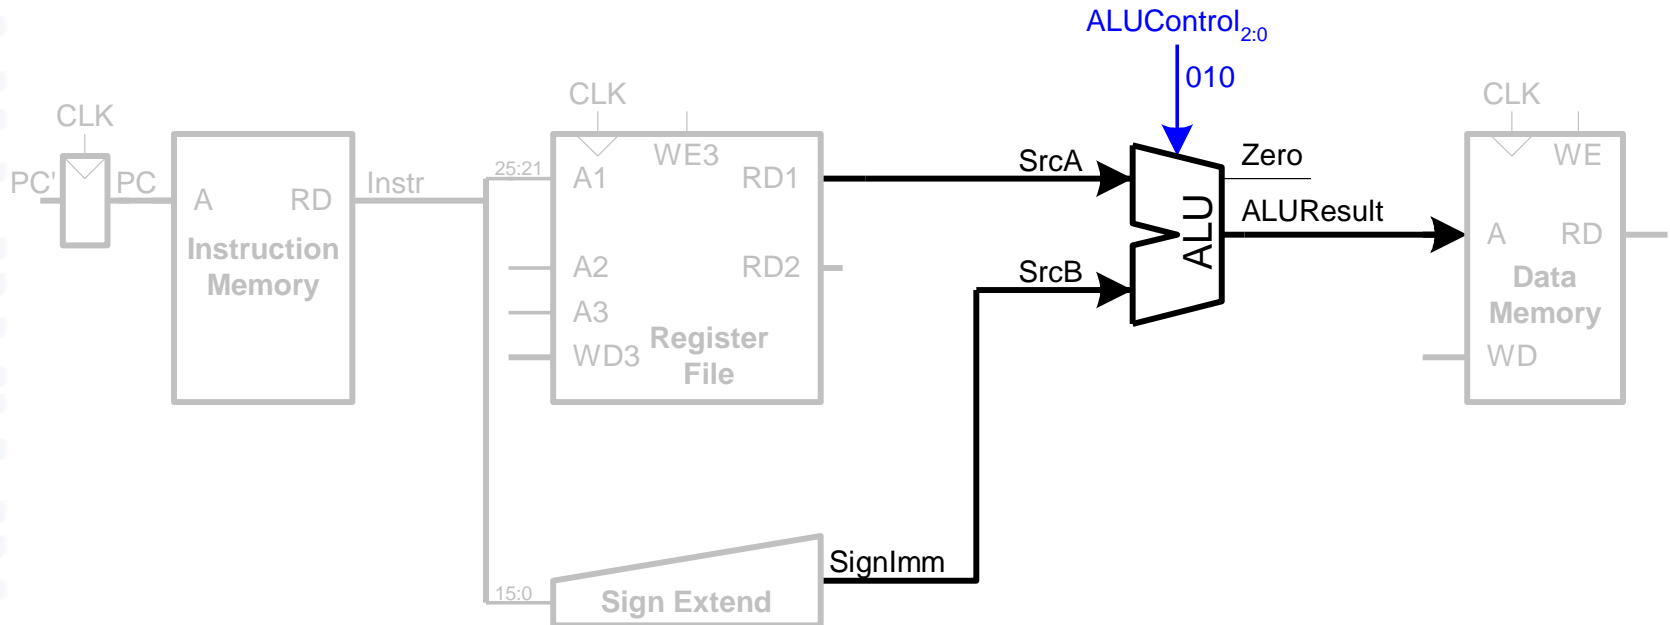
# Single-Cycle Datapath: $1_w$ Immediate

## STEP 3: Sign-extend the immediate $\text{SignExt}(\text{immed})$



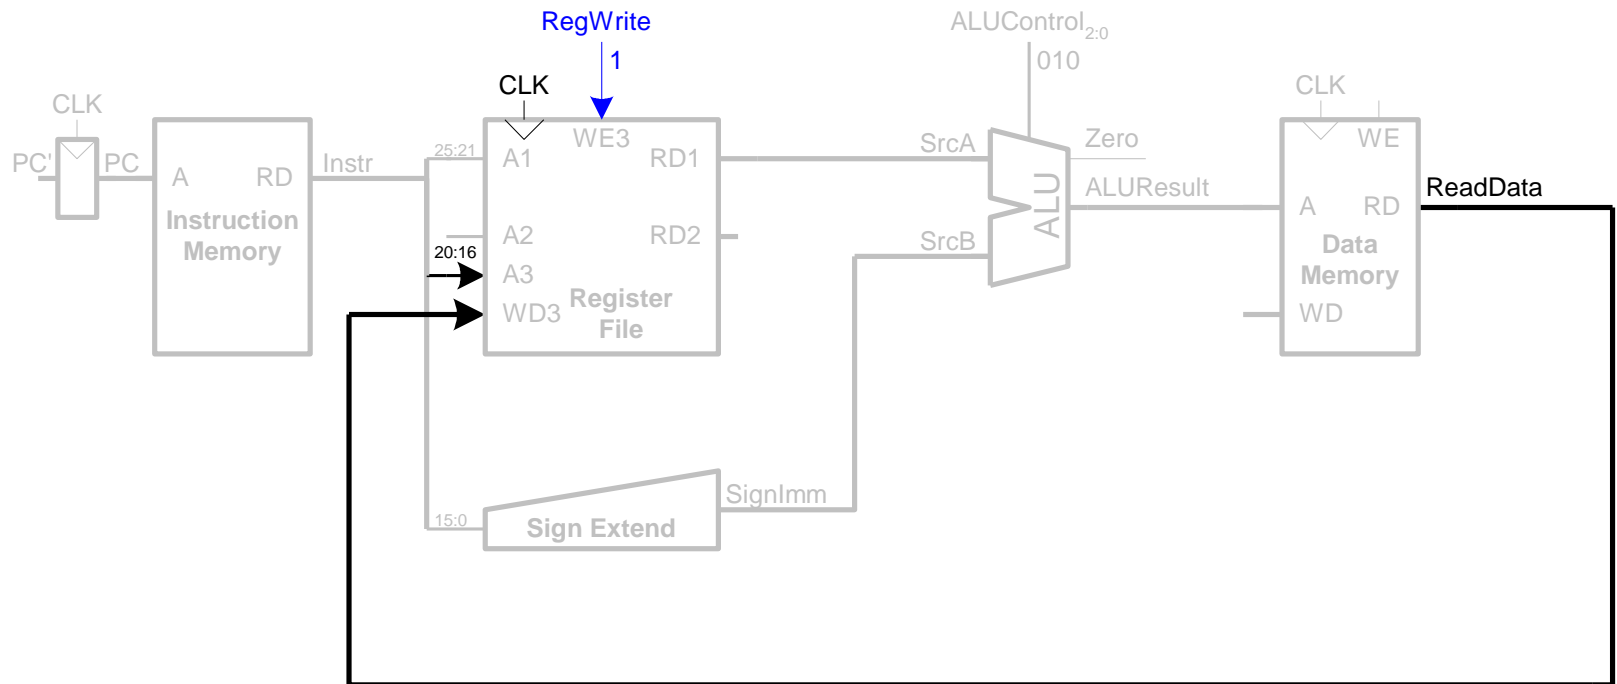
# Single-Cycle Datapath: lw address

**STEP 4:** Compute the memory address  
$$\text{addr} = \text{RF}[\text{rs}] + \text{SignExt}(\text{immed})$$



# Single-Cycle Datapath: $l_w$ Memory Read

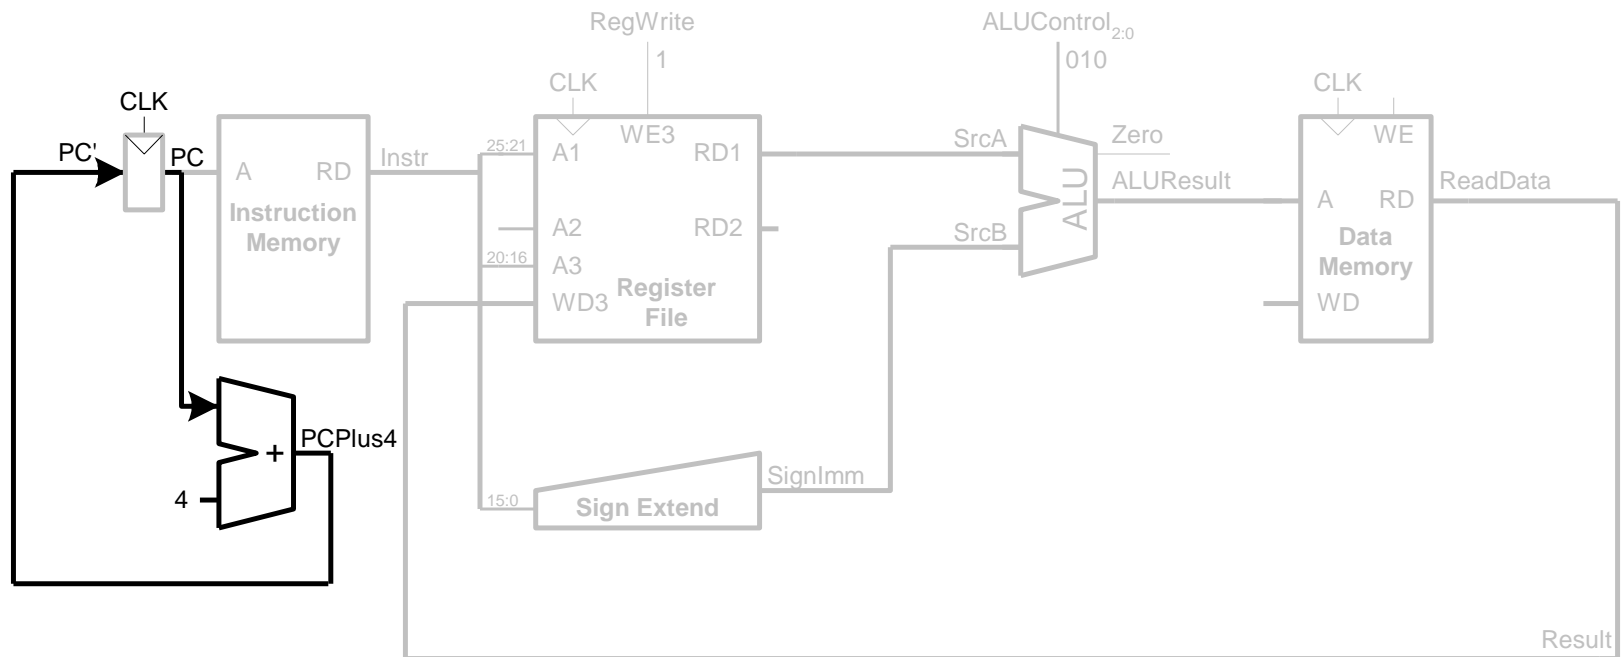
- STEP 5:** Read data from memory and write it back to register file:  $RF[rt] \leftarrow DM[addr]$



# Single-Cycle Datapath: $1_w$ PC Increment

**STEP 6:** Determine address of next instruction

$$PC \leftarrow PC + 4$$



# Full RTL Expression for $\perp_w$

IM[PC]

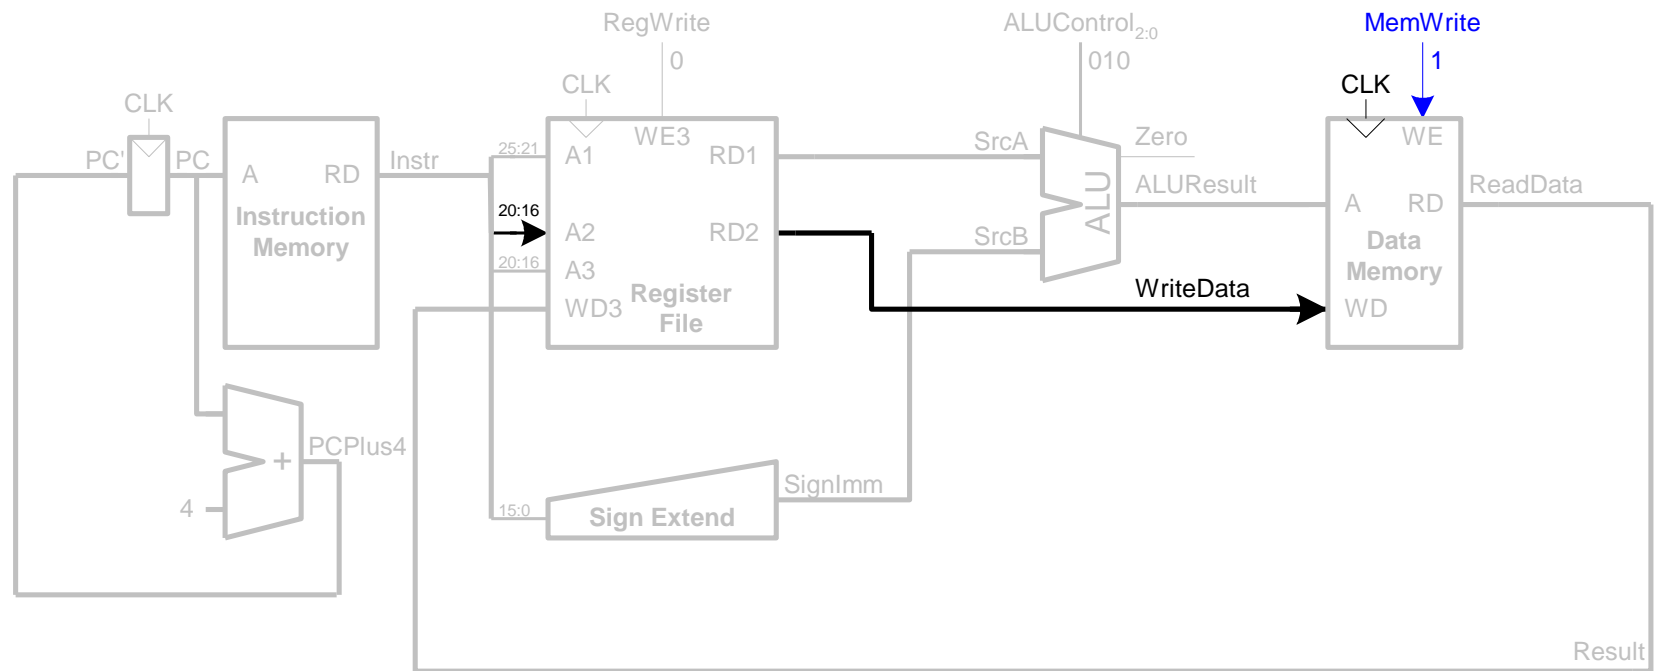
$RF[rt] \leftarrow DM[RF[rs] + \text{SignExt(immed)}]$

$PC \leftarrow PC + 4$



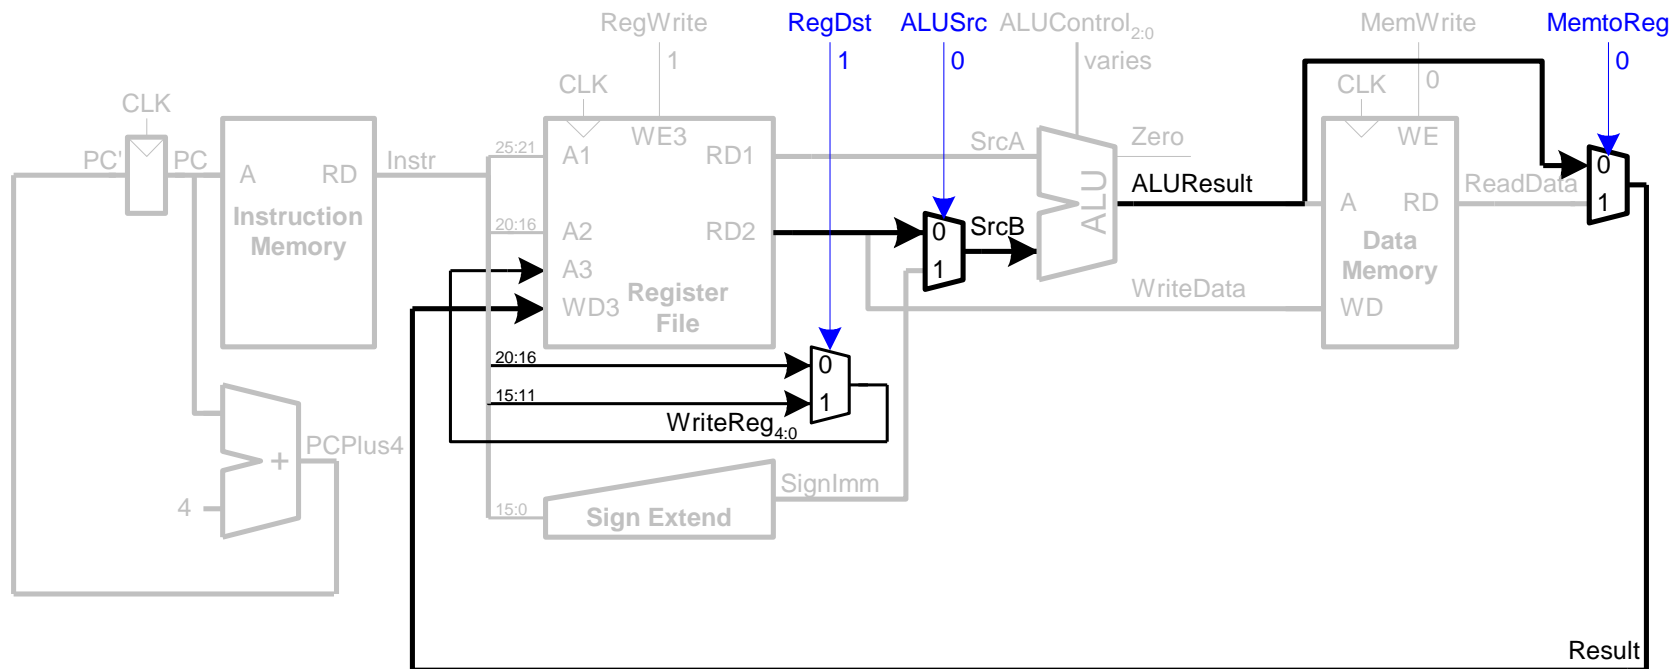
# Single-Cycle Datapath: sw

Write data in  $rt$  to memory:  $DM[addr] \leftarrow RF[rt]$



# Single-Cycle Datapath: R-Type


- Read from  $rs$  and  $rt$
- Write  $ALUResult$  to register file
- Write to  $rd$  (instead of  $rt$ )       $RF[rd] \leftarrow RF[rs] \text{ op } RF[rt]$



# Single-Cycle Datapath: beq

- Determine whether values in  $rs$  and  $rt$  are equal
- Calculate branch target address:

$$BTA = PC + 4 + \text{SignExt}(\text{immed}) \ll 2 \quad \# \ll 2 = 4x$$



# RTL Expression for beq

IM[PC]

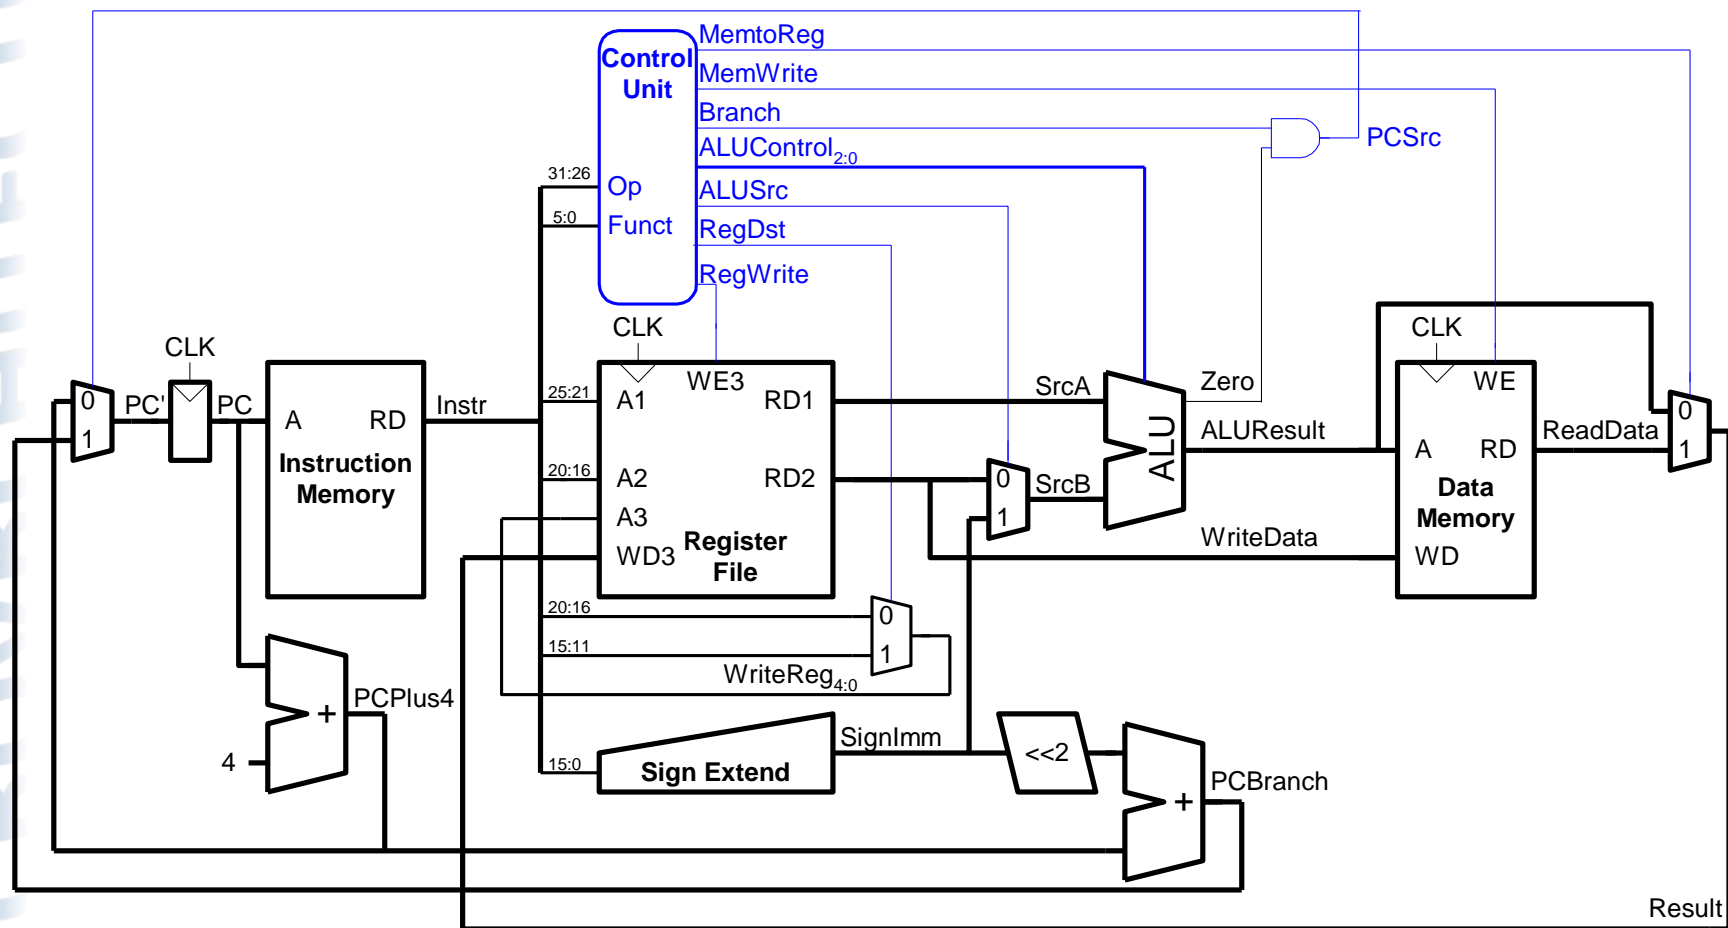
if (RF[rs] - RF[rt] == 0)

PC  $\leftarrow$  BTA

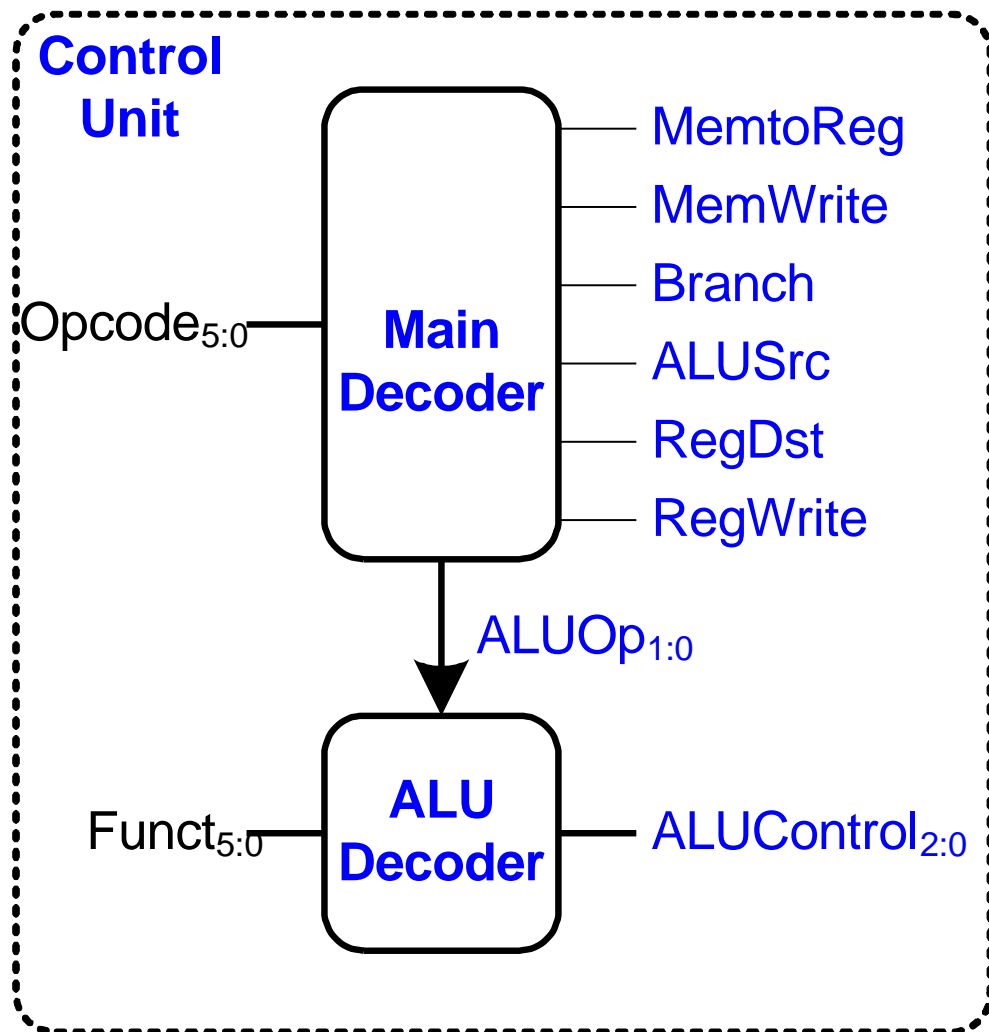
else

PC  $\leftarrow$  PC + 4


# Single-Cycle Processor



# Single-Cycle Control




# Review: ALU



$F_{2:0}$	Function
000	$A \& B$
001	$A   B$
010	$A + B$
011	not used
100	$A \& \sim B$
101	$A   \sim B$
110	$A - B$
111	SLT

# Review: ALU





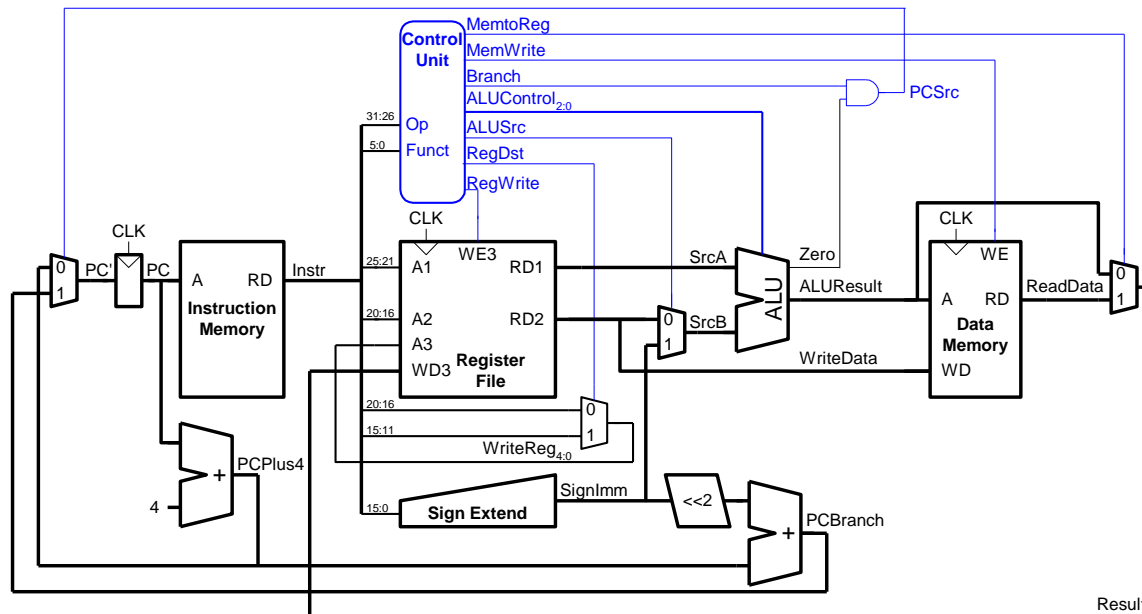
# Control Unit: ALU Decoder

ALUOp <sub>1:0</sub>	Meaning
00	Add (for lw, sw)
01	Subtract (for beq)
10	Look at funct (R-type)
11	Not Used

ALUOp <sub>1:0</sub>	funct	ALUControl <sub>2:0</sub>
00	X	010 (Add)
X1	X	110 (Subtract)
1X	100000 (add)	010 (Add)
1X	100010 (sub)	110 (Subtract)
1X	100100 (and)	000 (And)
1X	100101 (or)	001 (Or)
1X	101010 (slt)	111 (SLT)

# Control Unit: Main Decoder

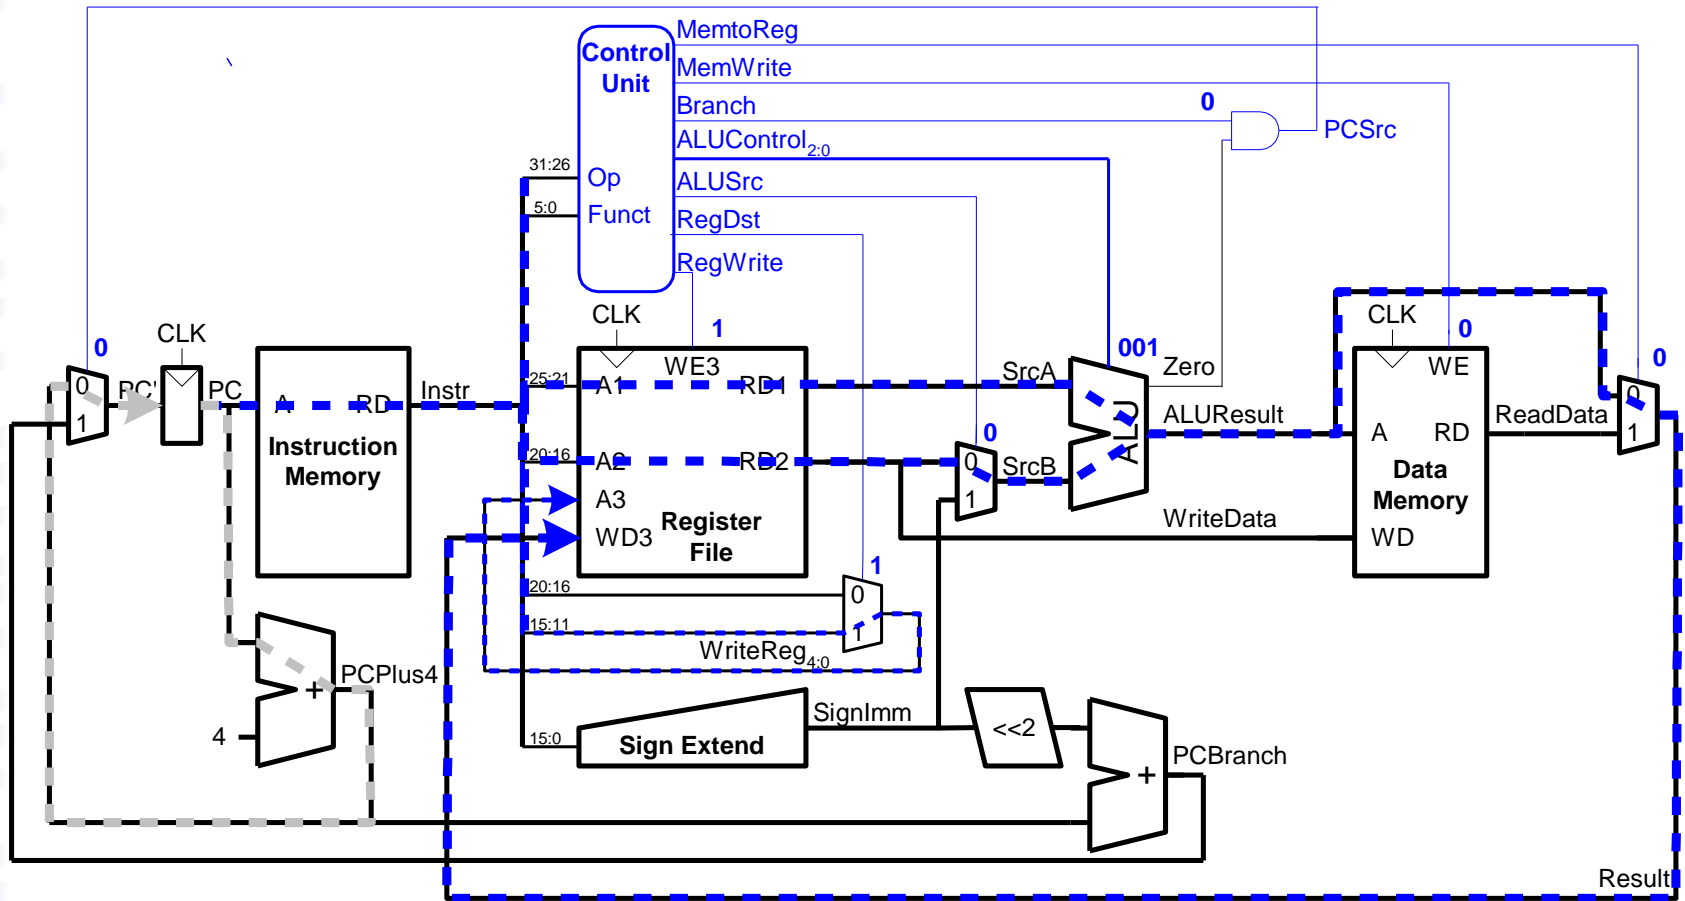
Instruction	Op <sub>5:0</sub>	RegWrite	RegDst	AluSrc	Branch	MemWrite	MemtoReg	ALUOp <sub>1:0</sub>
R-type	000000							
lw	100011							
sw	101011							
beq	000100							



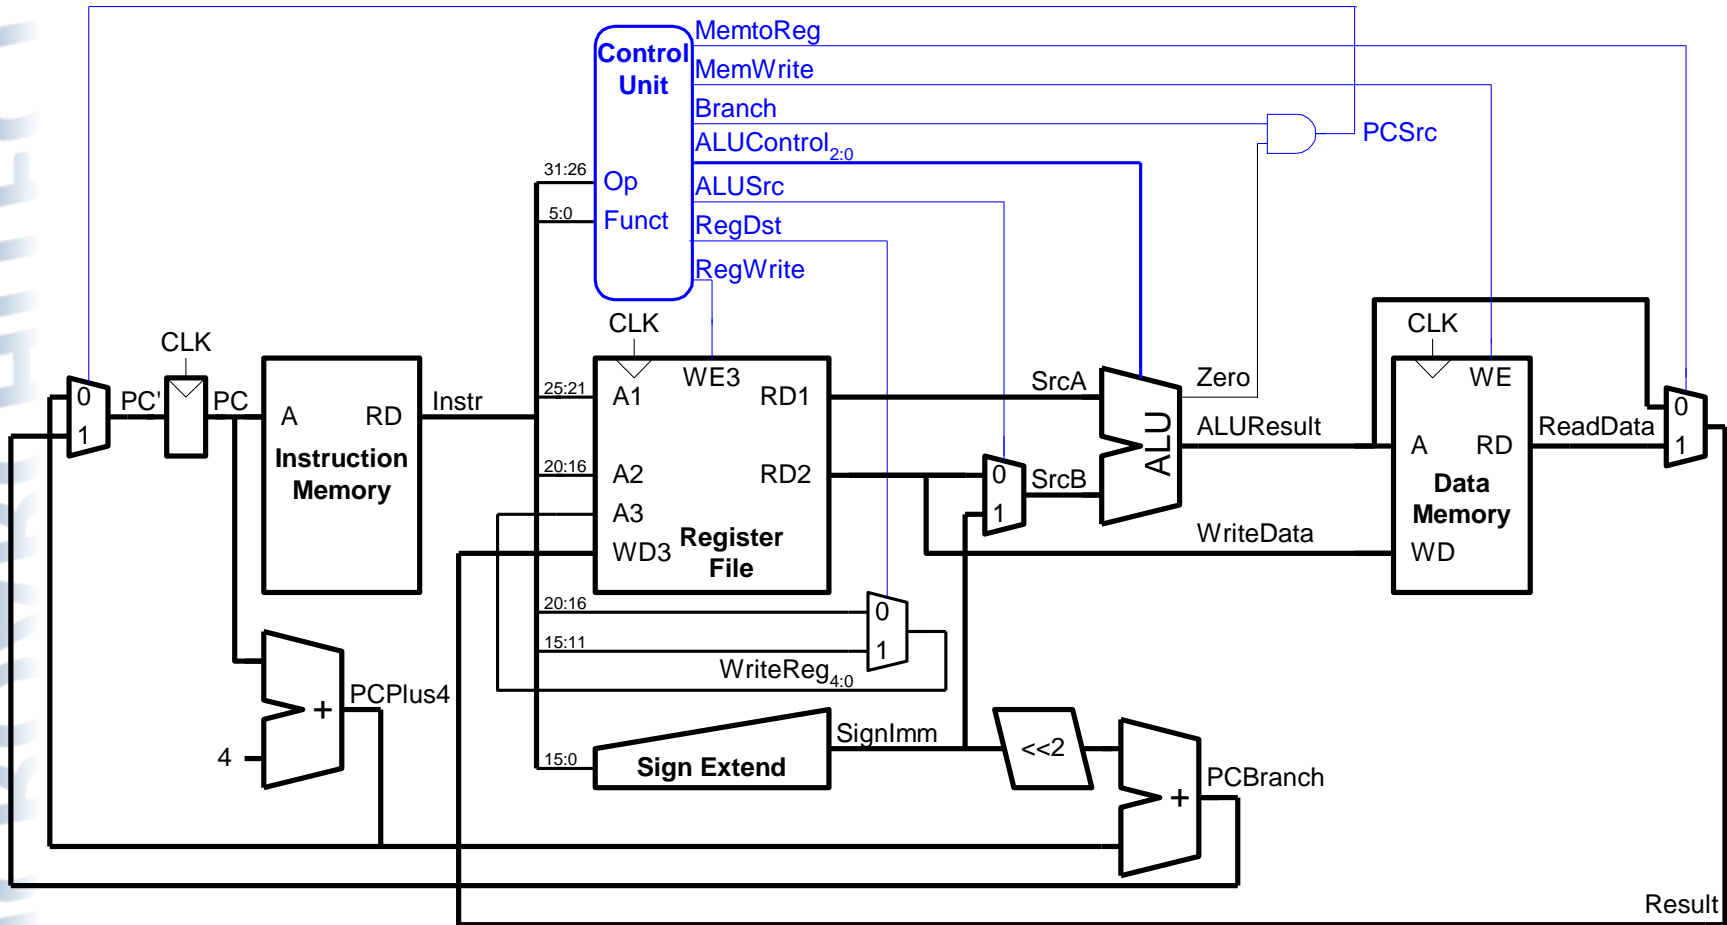
# Control Unit: Main Decoder

Instruction	Op <sub>5:0</sub>	RegWrite	RegDst	AluSrc	Branch	MemWrite	MemtoReg	ALUOp <sub>1:0</sub>
R-type	000000	1	1	0	0	0	0	10
lw	100011	1	0	1	0	0	1	00
sw	101011	0	X	1	0	1	X	00
beq	000100	0	X	0	1	0	X	01

# Single-Cycle Datapath: or



# Extended Functionality: addi



No change to datapath

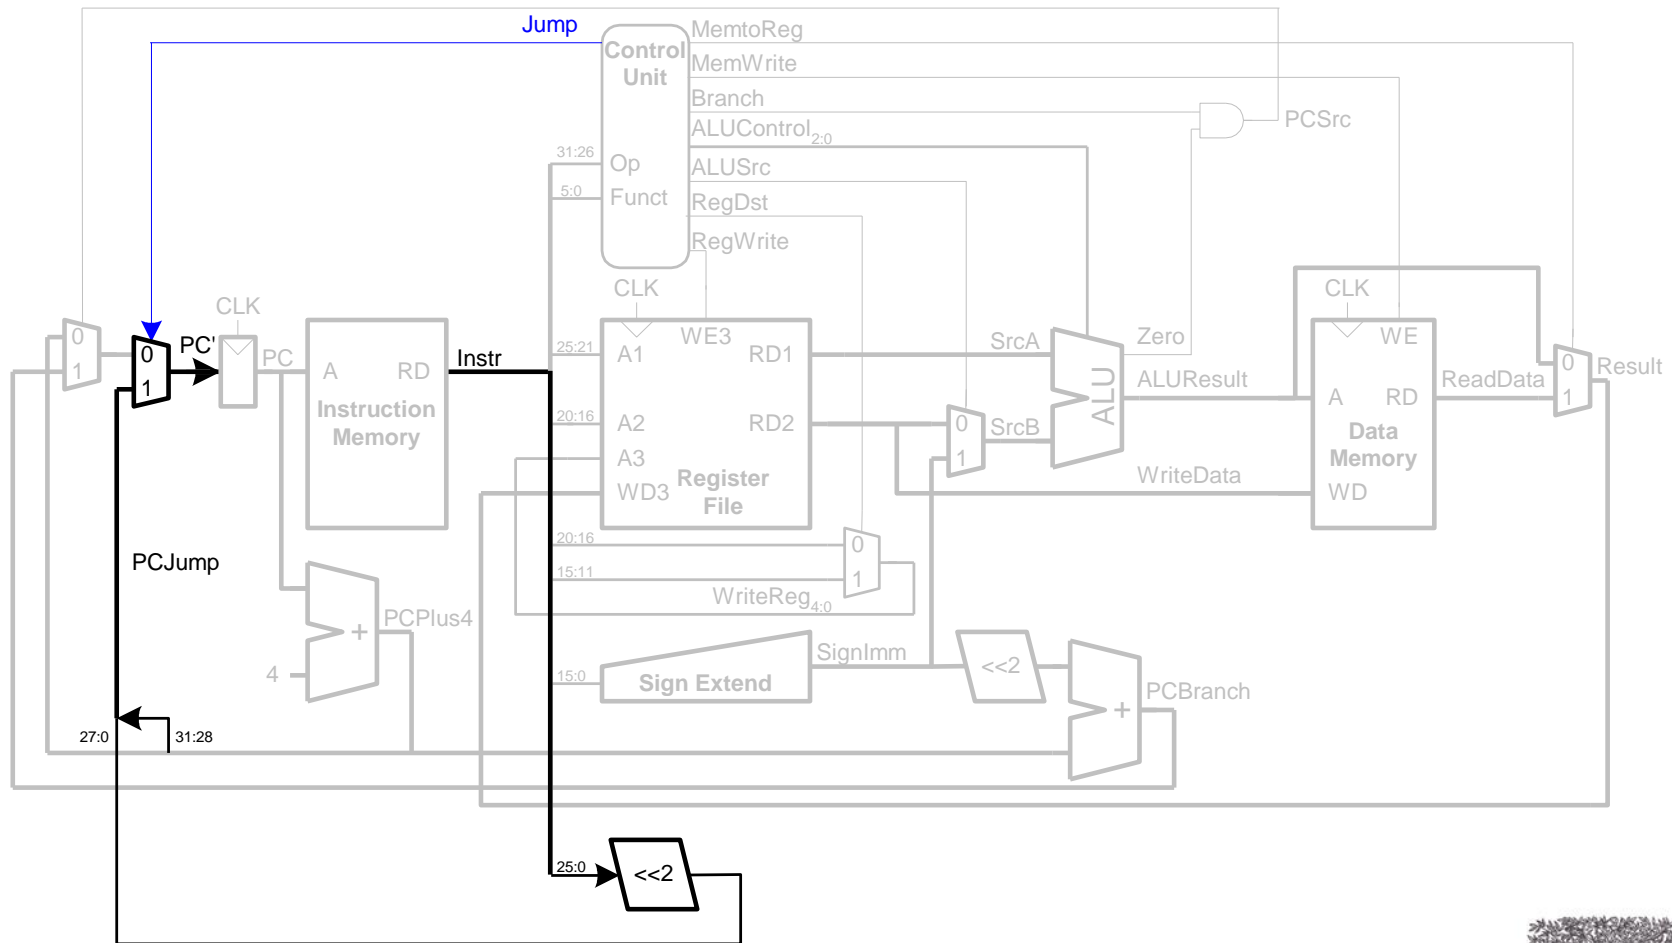
# Main Decoder table: addi

Instruction	Op <sub>5:0</sub>	RegWrite	RegDst	AluSrc	Branch	MemWrite	MemtoReg	ALUOp <sub>1:0</sub>
R-type	000000	1	1	0	0	0	0	10
lw	100011	1	0	1	0	0	1	00
sw	101011	0	X	1	0	1	X	00
beq	000100	0	X	0	1	0	X	01
<b>addi</b>	<b>001000</b>							

# Main Decoder table: addi

Instruction	Op <sub>5:0</sub>	RegWrite	RegDst	AluSrc	Branch	MemWrite	MemtoReg	ALUOp <sub>1:0</sub>
R-type	000000	1	1	0	0	0	0	10
lw	100011	1	0	1	0	0	1	00
sw	101011	0	X	1	0	1	X	00
beq	000100	0	X	0	1	0	X	01
<b>addi</b>	<b>001000</b>	<b>1</b>	<b>0</b>	<b>1</b>	<b>0</b>	<b>0</b>	<b>0</b>	<b>00</b>

# Extended Functionality: j





# Main Decoder table: j

Instruction	Op <sub>5:0</sub>	RegWrite	RegDst	AluSrc	Branch	MemWrite	MemtoReg	ALUOp <sub>1:0</sub>	Jump
R-type	000000	1	1	0	0	0	0	10	0
lw	100011	1	0	1	0	0	1	00	0
sw	101011	0	X	1	0	1	X	00	0
beq	000100	0	X	0	1	0	X	01	0
j	000010								

# Main Decoder table: j

Instruction	Op <sub>5:0</sub>	RegWrite	RegDst	AluSrc	Branch	MemWrite	MemtoReg	ALUOp <sub>1:0</sub>	Jump
R-type	000000	1	1	0	0	0	0	10	0
lw	100011	1	0	1	0	0	1	00	0
sw	101011	0	X	1	0	1	X	00	0
beq	000100	0	X	0	1	0	X	01	0
j	000010	0	X	X	X	0	X	XX	1

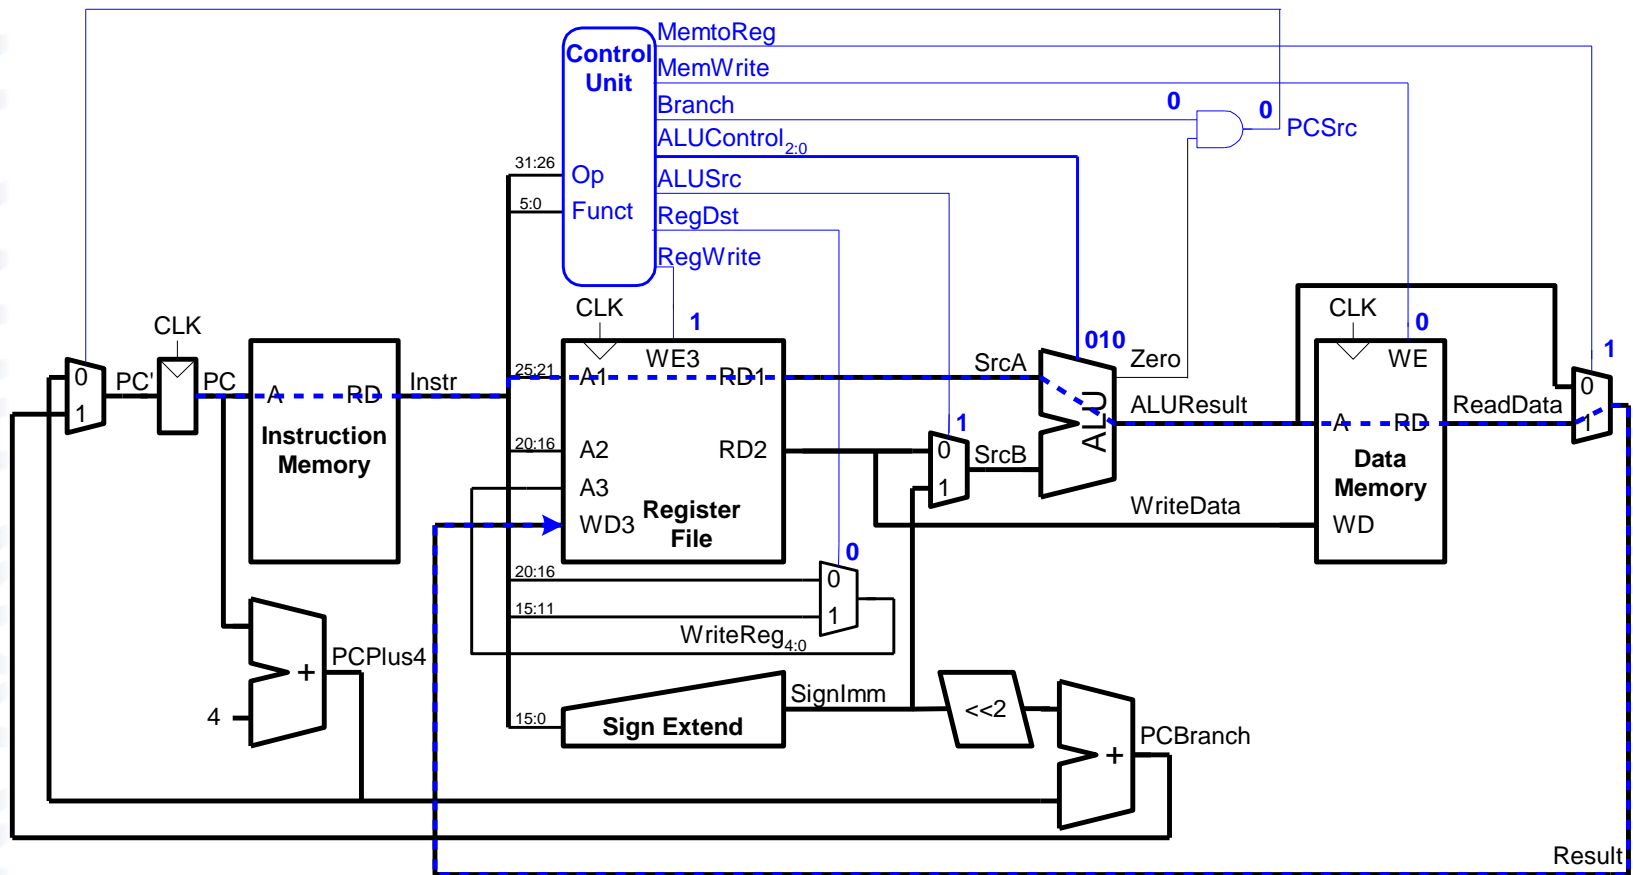
# Review: Processor Performance

## Program Execution Time

$$= (\text{\#instructions})(\text{cycles/instruction})(\text{seconds/cycle})$$

$$= IC \times CPI \times T_c$$

# Single-Cycle Performance



$T_C$  limited by critical path ( $1w$ )

# Single-Cycle Performance

- Single-cycle critical path:

$$T_c = t_{pcq\_PC} + t_{mem} + \max(t_{RFread}, t_{sext} + t_{mux}) + t_{ALU} + t_{mem} + t_{mux} + t_{RFsetup}$$

- Typically, limiting paths are:

- memory, ALU, register file

- $T_c = t_{pcq\_PC} + 2t_{mem} + t_{RFread} + t_{mux} + t_{ALU} + t_{RFsetup}$



# Single-Cycle Performance Example

Element	Parameter	Delay (ps)
Register clock-to-Q	$t_{pcq\_PC}$	30
Register setup	$t_{setup}$	20
Multiplexer	$t_{mux}$	25
ALU	$t_{ALU}$	200
Memory read	$t_{mem}$	250
Register file read	$t_{RFread}$	150
Register file setup	$t_{RFsetup}$	20

$$T_c = ?$$

# Single-Cycle Performance Example

Element	Parameter	Delay (ps)
Register clock-to-Q	$t_{pcq\_PC}$	30
Register setup	$t_{setup}$	20
Multiplexer	$t_{mux}$	25
ALU	$t_{ALU}$	200
Memory read	$t_{mem}$	250
Register file read	$t_{RFread}$	150
Register file setup	$t_{RFsetup}$	20

$$\begin{aligned}T_c &= t_{pcq\_PC} + 2t_{mem} + t_{RFread} + t_{mux} + t_{ALU} + t_{RFsetup} \\&= [30 + 2(250) + 150 + 25 + 200 + 20] \text{ ps} \\&= 925 \text{ ps} \quad [f_{clk} = 1/0.925 \text{ GHz} = 1.08 \text{ GHz}]\end{aligned}$$

# Single-Cycle Performance Example

Program with IC = 100 billion instructions:

$$\begin{aligned}\text{Execution Time} &= \text{IC} \times \text{CPI} \times T_C \\ &= (100 \times 10^9)(1)(925 \times 10^{-12} \text{ s}) \\ &= \mathbf{92.5 \text{ seconds}}\end{aligned}$$

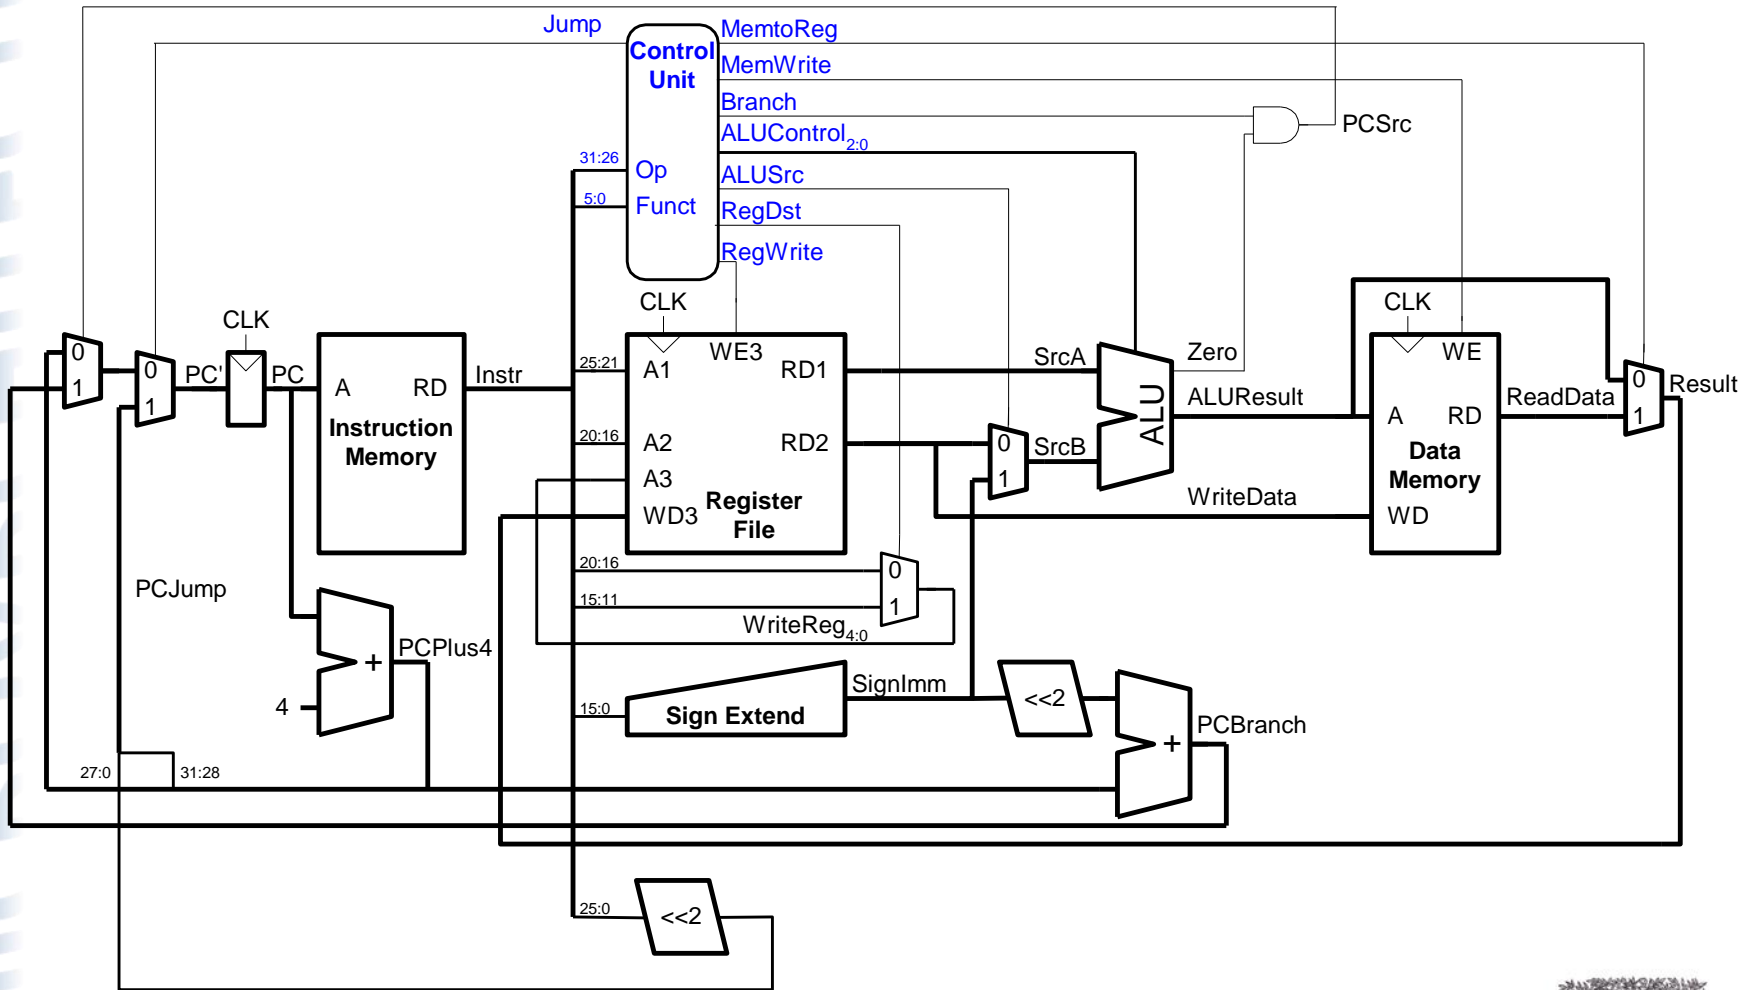


# Evaluation of Single-Cycle Processor

## Pros and cons of single-cycle implementation:

- + simple design
- + 1 cycle per every instruction
- slow cycle time
  - limited by longest instruction ( $1_w$ )
- HW: 2 adders + ALU; 2 memories

# Review: Single-Cycle Processor



# System Verilog Model

```
//-----  
// by David Harris, in Chapter 7 of DDCA 2nd ed. textbook  
// Top level system of Fig 7.59 including MIPS and memories  
//-----  
  
module top (input  logic      clk, reset,  
            output logic[31:0] writedata, dataadr,  
            output logic      memwrite);  
  
    logic [31:0] pc, instr, readdata;  
  
    // instantiate processor and memories  
    mips mips (clk, reset, pc, instr, memwrite, dataadr, writedata, readdata);  
    imem imem (pc[7:2], instr);  
    dmem dmem (clk, memwrite, dataadr, writedata, readdata);  
  
endmodule
```

# Sys. Verilog Model of Data Memory

```
//-----  
// External data memory used by MIPS single-cycle processor  
//-----  
module dmem (input logic      clk, we,  
              input logic[31:0] a, wd,  
              output logic[31:0] rd);  
  
  logic [31:0] RAM[63:0];  
  
  assign rd = RAM[a[31:2]]; // word-aligned read  
  
  always_ff @(posedge clk)  
    if (we)  
      RAM[a[31:2]] <= wd; // word-aligned write  
endmodule
```

# Sys. Verilog Model of Instr. Memory

```
module imem (input logic[5:0] addr, output reg [31:0] instr);

// imem is modeled as a lookup table, a stored-program byte-addressable ROM
always_comb
    case ({addr,2'b00})
//          address          instruction
//          -----          -
        8'h00: instr = 32'h20020005;
        8'h04: instr = 32'h2003000c;
        8'h08: instr = 32'h2067fff7;
        8'h0c: instr = 32'h00e22025;
        8'h10: instr = 32'h00642824;
        8'h14: instr = 32'h00a42820;
        8'h18: instr = 32'h10a7000a;
        8'h1c: instr = 32'h0064202a;
        8'h20: instr = 32'h10800001;
        default: instr = {32{1'bx}}; // unknown instruction
    endcase
endmodule
```



# Alternate Model of Instr. Memory

```
module imem (input  logic[5:0] addr,
             output logic[31:0] instr);

logic [31:0] RAM[63:0];

// imem is RAM, loaded from memfile.dat file with hex values at startup
initial
begin
    $readmemh("memfile.dat", RAM);
end

assign instr = RAM[addr]; // instr at RAM[addr] is read out

endmodule

// imem can be created with CoreGen for Xilinx synthesis
```

# Sys. Verilog Model of MIPS processor

```
// single-cycle MIPS processor
module mips (input logic      clk, reset,
             output logic[31:0] pc,
             input logic[31:0] instr,
             output logic      memwrite,
             output logic[31:0] aluout, writedata,
             input logic[31:0] readdata);

    logic      memtoreg, pcsrc, zero, alusrc, regdst, regwrite, jump;
    logic [2:0] alucontrol;

    controller c (instr[31:26], instr[5:0], zero, memtoreg, memwrite, pcsrc,
                  alusrc, regdst, regwrite, jump, alucontrol);

    datapath dp (clk, reset, memtoreg, pcsrc, alusrc, regdst, regwrite, jump,
                 alucontrol, zero, pc, instr, aluout, writedata, readdata);

endmodule
```



# Sys. Verilog Model of Controller

```
module controller ( input logic[5:0] op, funct,
                    input logic      zero,
                    output logic      memtoreg, memwrite,
                    output logic      pcsrc, alusrc,
                    output logic      regdst, regwrite,
                    output logic      jump,
                    output logic[2:0] alucontrol);

    logic [1:0] aluop;
    logic      branch;

    maindec md (op, memtoreg, memwrite, branch, alusrc, regdst, regwrite,
                jump, aluop);

    aludec ad (funct, aluop, alucontrol);

    assign pcsrc = branch & zero;

endmodule
```



# Sys. Verilog Model of Main Decoder

```
module maindec (input  logic[5:0] op,  
                output logic memtoreg, memwrite, branch,  
                output logic alusrc, regdst, regwrite, jump,  
                output logic[1:0] aluop );  
  
    logic [8:0] controls;  
  
    assign {regwrite, regdst, alusrc, branch, memwrite,  
            memtoreg, aluop, jump} = controls;  
  
    always_comb  
    case(op)  
        6'b000000: controls = 9'b1100000010; //Rtype  
        6'b100011: controls = 9'b101001000; //LW  
        6'b101011: controls = 9'b001010000; //SW  
        6'b000100: controls = 9'b000100001; //BEQ  
        6'b001000: controls = 9'b101000000; //ADDI  
        6'b000010: controls = 9'b0000000100; //J  
        default:   controls = 9'bxxxxxxxx; //illegal op  
    endcase  
endmodule
```



# Sys. Verilog Model of ALU Decoder

```
module aludec (input  logic[5:0] funct,
               input  logic[1:0] aluop,
               output logic [2:0] alucontrol);

always_comb
case(aluop)
  2'b00: alucontrol = 3'b010; // add (for lw/sw/addi)
  2'b01: alucontrol = 3'b110; // sub (for beq)
  default: case(funct)          // R-TYPE instructions
    6'b100000: alucontrol = 3'b010; // ADD
    6'b100010: alucontrol = 3'b110; // SUB
    6'b100100: alucontrol = 3'b000; // AND
    6'b100101: alucontrol = 3'b001; // OR
    6'b101010: alucontrol = 3'b111; // SLT
    default:   alucontrol = 3'bxxx; // ???
  endcase
endcase
endmodule
```



# Sys. Verilog Model of Datapath

```
module datapath (input logic clk, reset, memtoreg, pcsrc, alusrc, regdst,  
                input logic regwrite, jump, input logic[2:0] alucontrol,  
                output logic zero, output logic[31:0] pc,  
                input logic[31:0] instr,  
                output logic[31:0] aluout, writedata,  
                input logic[31:0] readdata);  
  
    logic [4:0] writereg;  
    logic [31:0] pcnext, pcnextbr, pcplus4, pcbranch;  
    logic [31:0] signimm, signimmsh, srca, srcb, result;  
  
    // next PC logic  
    flopr #(32) pcreg(clk, reset, pcnext, pc);  
    adder      pcadd1(pc, 32'b100, pcplus4);  
    sl2        immsh(signimm, signimmsh);  
    adder      pcadd2(pcplus4, signimmsh, pcbranch);  
    mux2 #(32) pcbrmux(pcplus4, pcbranch, pcsrc,  
                      pcnextbr);  
    mux2 #(32) pcmux(pcnextbr, {pcplus4[31:28],  
                          instr[25:0], 2'b00}, jump, pcnext);
```



# Sys. Verilog Model of Datapath

```
// register file logic
```

```
regfile  rf (clk, regwrite, instr[25:21], instr[20:16], writereg,  
            result, srca, writedata);
```

```
mux2 #(5)  wrmux (instr[20:16], instr[15:11], regdst, writereg);
```

```
mux2 #(32) resmux (aluout, readdata, memtoreg, result);
```

```
signext    se (instr[15:0], signimm);
```

```
// ALU logic
```

```
mux2 #(32) srcbmux (writedata, signimm, alusrc, srcb);
```

```
alu        alu (srca, srcb, alucontrol, aluout, zero);
```

```
endmodule
```

# Sys. Verilog Model of Register File

```
module regfile (input  logic clk, we3,
                input  logic[4:0] ra1, ra2, wa3,
                input  logic[31:0] wd3,
                output logic[31:0] rd1, rd2);

    logic [31:0] rf [31:0];

    // three ported register file: read two ports combinationaly
    // write third port on rising edge of clock. Register0 hardwired to 0.

    always_ff
        if (we3)
            rf [wa3] <= wd3;

    assign rd1 = (ra1 != 0) ? rf [ra1] : 0;
    assign rd2 = (ra2 != 0) ? rf[ ra2] : 0;

endmodule
```



# Sys. Verilog Models of Other Parts

```
module adder (input  logic[31:0] a, b,  
              output logic[31:0] y);  
    assign y = a + b;  
endmodule
```

```
module sl2 (input  logic[31:0] a,  
            output logic[31:0] y);  
    // shift left by 2  
    assign y = {a[29:0], 2'b00};  
endmodule
```

```
module signext (input  logic[15:0] a,  
                output logic[31:0] y);  
  
    assign y = {{16{a[15]}}, a};  
endmodule
```

# Sys. Verilog for Parameterized Parts

```
module flopr #(parameter WIDTH = 8)
    (input logic clk, reset,
     input logic[WIDTH-1:0] d,
     output logic[WIDTH-1:0] q);
```

```
    always_ff@(posedge clk, posedge reset)
        if (reset) q <= 0;
        else      q <= d;
endmodule
```

```
module mux2 #(parameter WIDTH = 8)
    (input  logic[WIDTH-1:0] d0, d1,
     input  logic s,
     output logic[WIDTH-1:0] y);
    assign y = s ? d1 : d0;
endmodule
```



# Review: Exceptions

- Unscheduled function call to *exception handler*
- Caused by:
  - Hardware, also called an *interrupt*, e.g. keyboard
  - Software, also called *traps*, e.g. undefined instruction
- When exception occurs, the processor:
  - Records cause of exception (Cause register)
  - Jumps to exception handler (0x80000180)
  - Returns to program (EPC register)



# Example Exception

sequential circuits.]

Can we design a spiff

Figure 2.1 shows a inputs, A and B, and on box indicates that it is this case, the function is

KeyAccess



The network KeyServer, which is required by KeyServer controlled programs, cannot grant you permission to run this program. If you think you have received this message in error, please contact your KeyServer Administrator.

Visio.exe - Application Error



The exception unknown software exception (0xc06d007e) occurred in the application at location 0x7c81eb33.

OK

words, we say the output Y is a function of the two inputs A and B where the function performed is A OR B.]

The *implementation* of the combinational circuit is independent of its functionality. Figure 2.1 and Figure 2.2 show two possible implementa-



# Review: Exception Registers

- Not part of register file; in Coprocessor 0
  - Cause
    - Records cause of exception
    - Coprocessor 0 register 13
  - EPC (Exception PC)
    - Records PC where exception occurred
    - Coprocessor 0 register 14
- Move from Coprocessor 0
  - `mfc0 $t0, Cause (=mfc0 $t0, $13)`
  - Moves contents of Cause into \$t0

**mfc0**

010000	00000	\$t0 (8)	Cause (13)	000000000000
31:26	25:21	20:16	15:11	10:0

# Review: Exception Causes

Exception	Cause
Hardware Interrupt	0x00000000
System Call	0x00000020
Breakpoint / Divide by 0	0x00000024
Undefined Instruction	0x00000028
Arithmetic Overflow	0x00000030

**Extend single-cycle MIPS processor to handle last two types of exceptions**

# Exception RTLs

## Undefined Instruction

IM[PC]

... # problem in decoding (bad op or func)

Cause  $\leftarrow 40$  # = 0x28

EPC  $\leftarrow$  PC

PC  $\leftarrow$  0x80000180 #Exception handler address

## Arithmetic Overflow

IM[PC]

... # ALU operation overflows

Cause  $\leftarrow 48$  # = 0x30

EPC  $\leftarrow$  PC

PC  $\leftarrow$  0x80000180 #Exception handler address

**mfc0 instruction** (e.g. mfc0 \$t1, \$13)

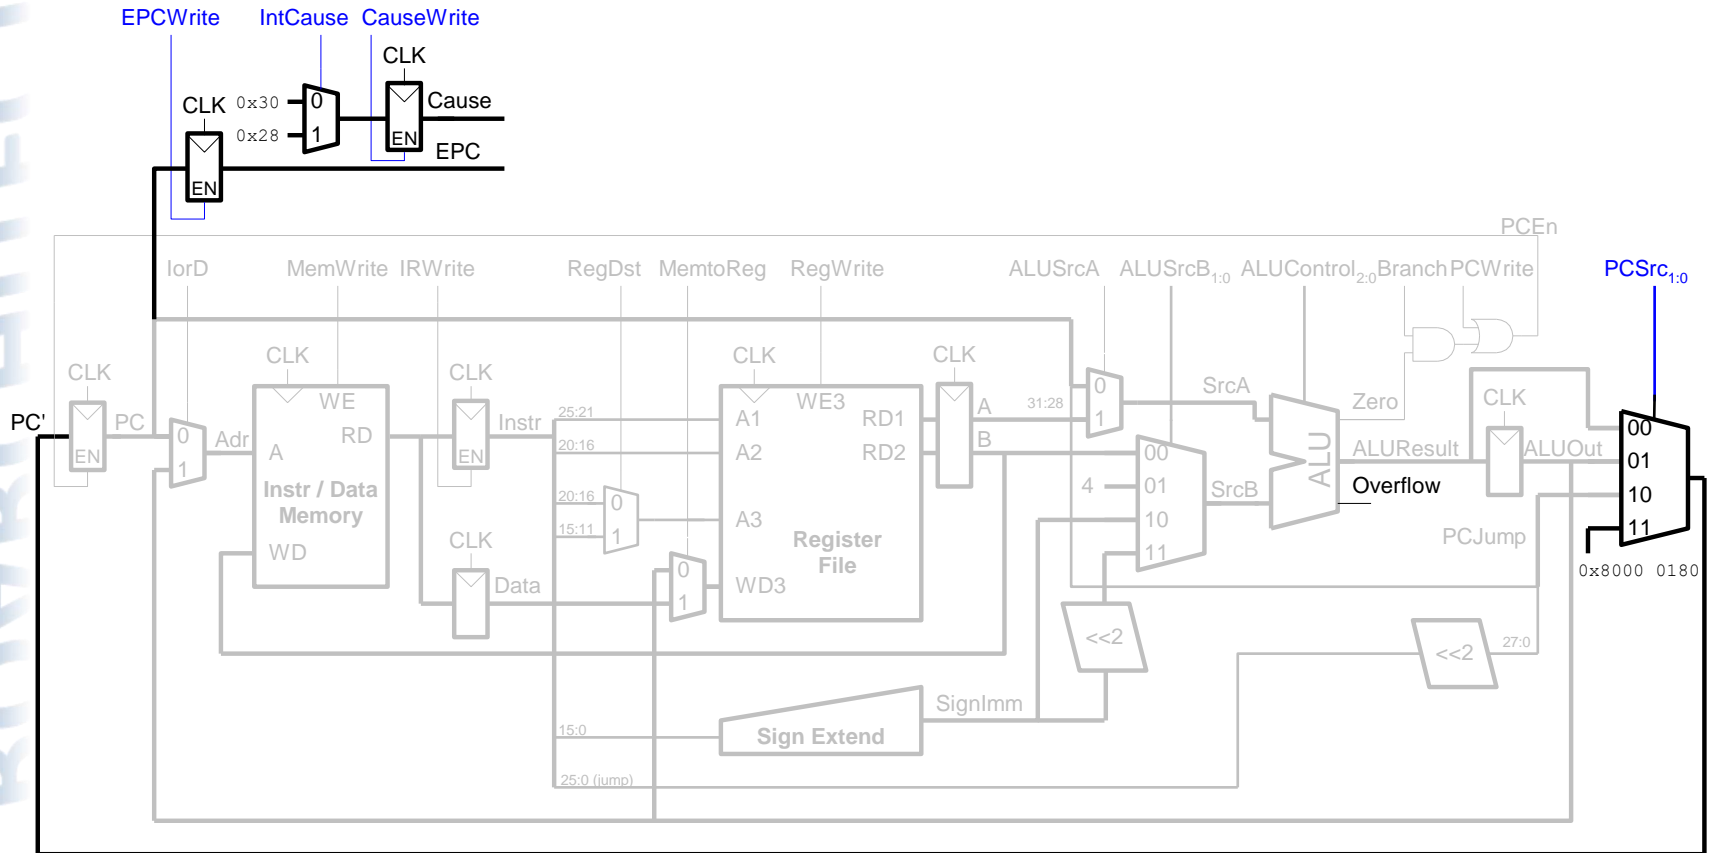
IM[PC]

RF[rt]  $\leftarrow$  RFc0[rd]

PC  $\leftarrow$  PC + 4




# Exception Hardware: EPC & Cause



Never mind the *multi-cycle* datapath, focus on the exception hardware.



# Exception Hardware: mfc0



Never mind the *multi-cycle* datapath, focus on the exception hardware.

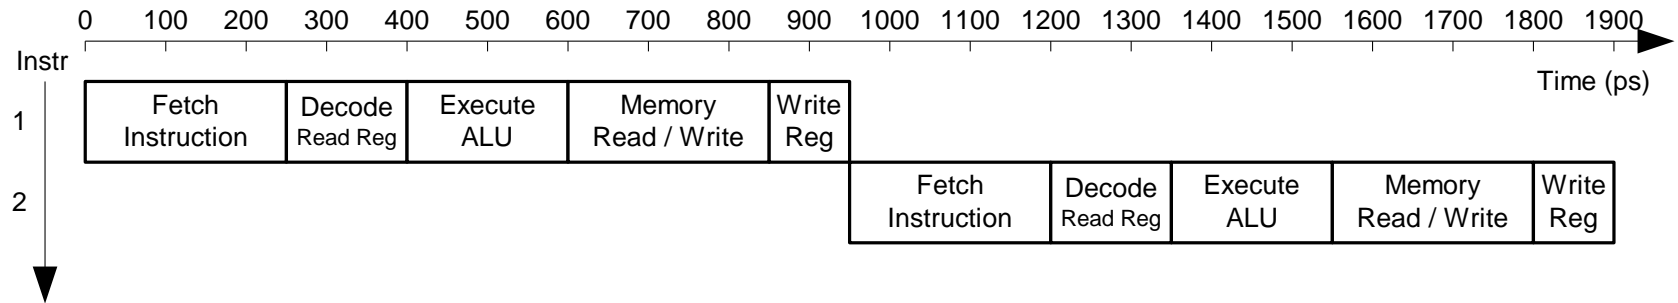


# Pipelined MIPS Processor

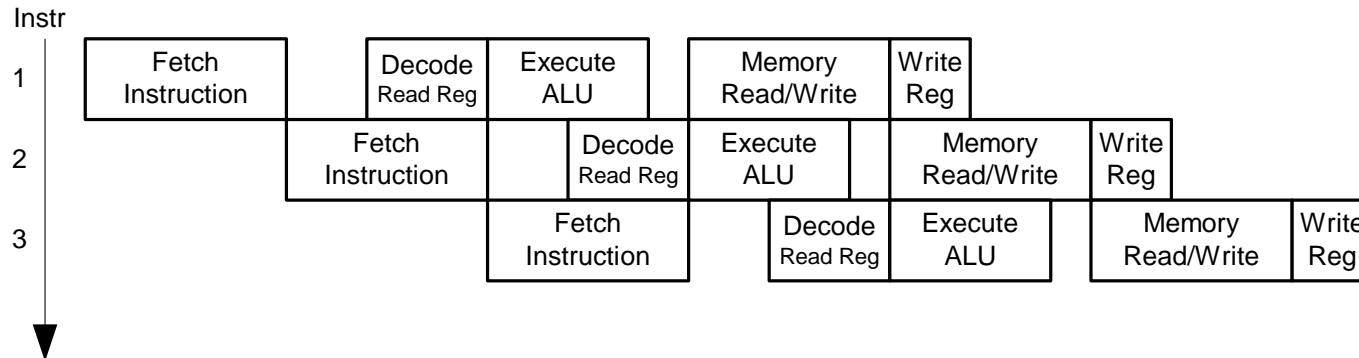
- Temporal parallelism
- Divide single-cycle processor into 5 stages:
  - Fetch
  - Decode
  - Execute
  - Memory
  - Writeback
- Add pipeline registers between stages

# Single-Cycle vs. Pipelined

## Single-Cycle

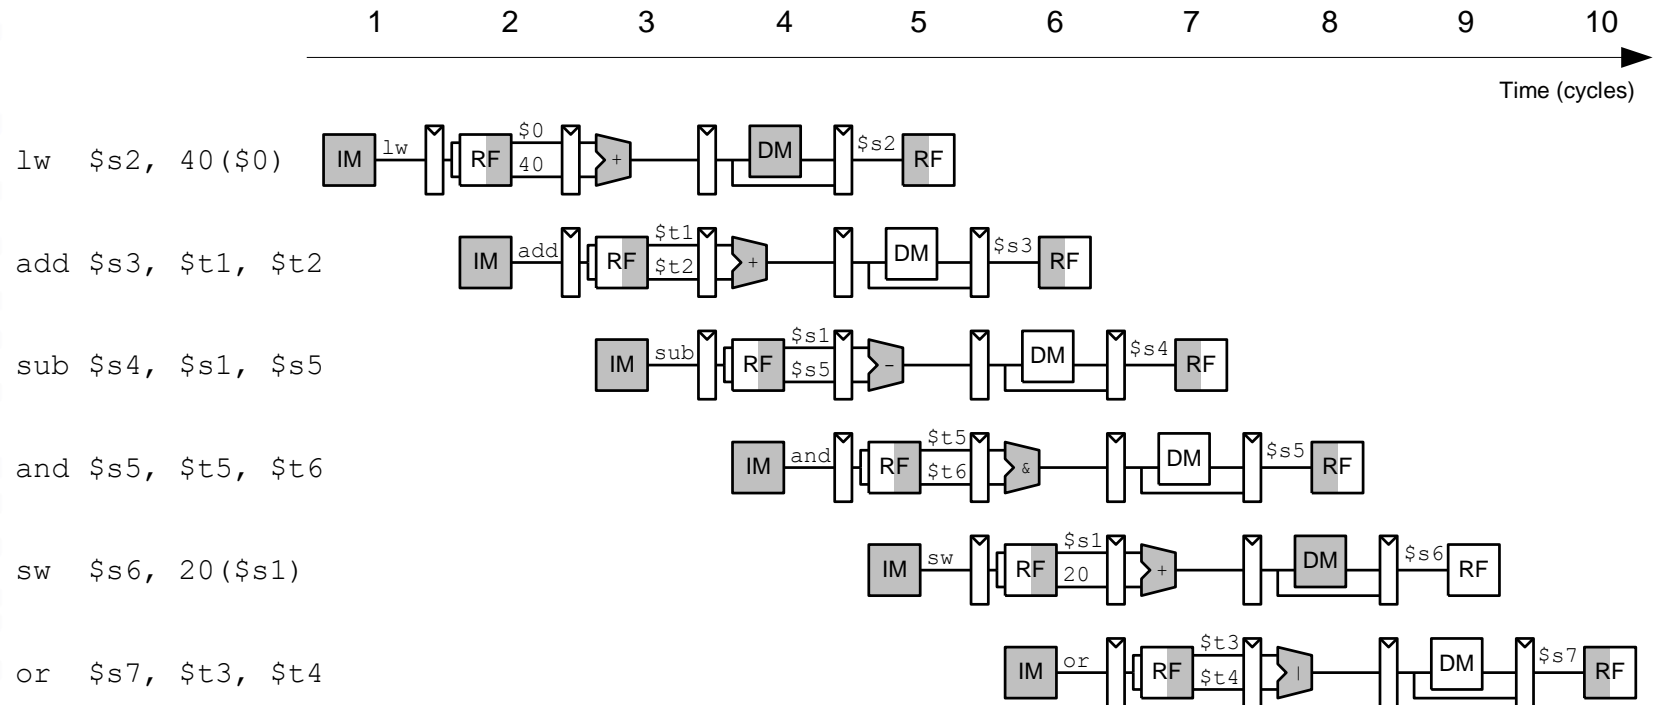


## Pipelined

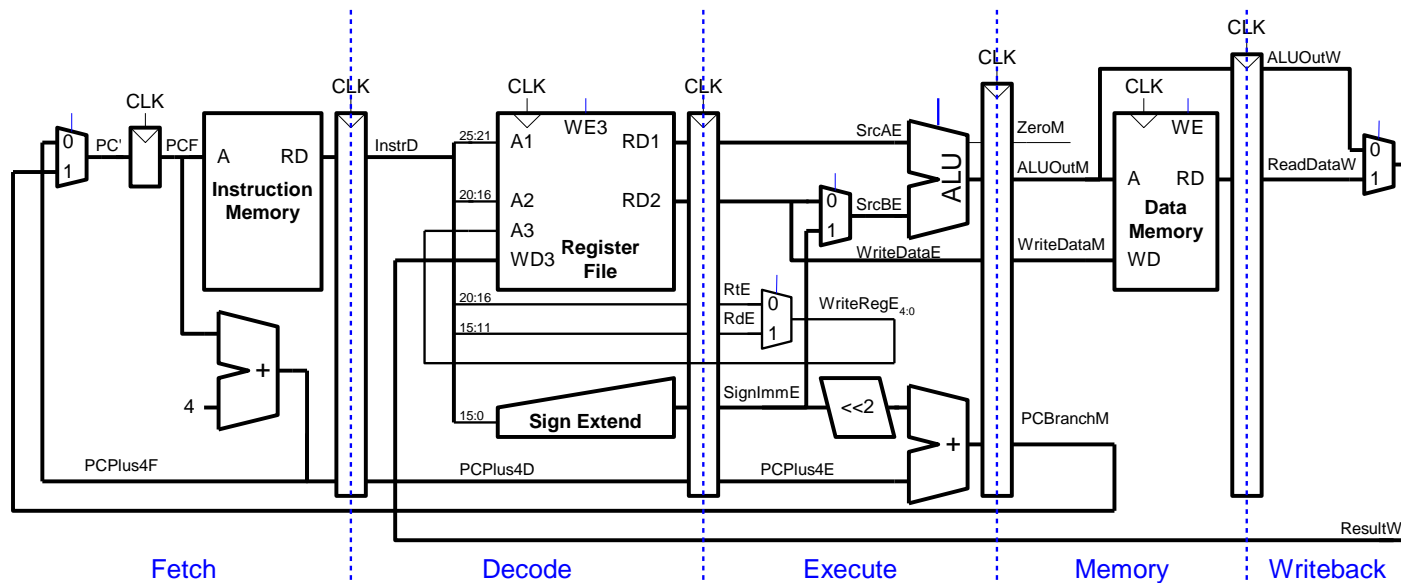
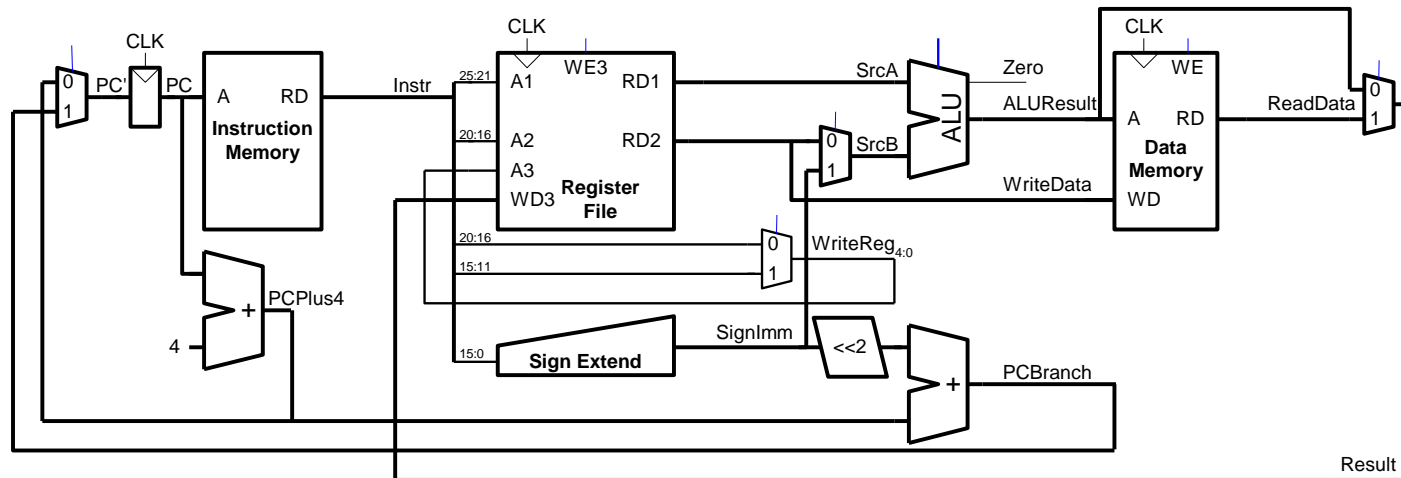




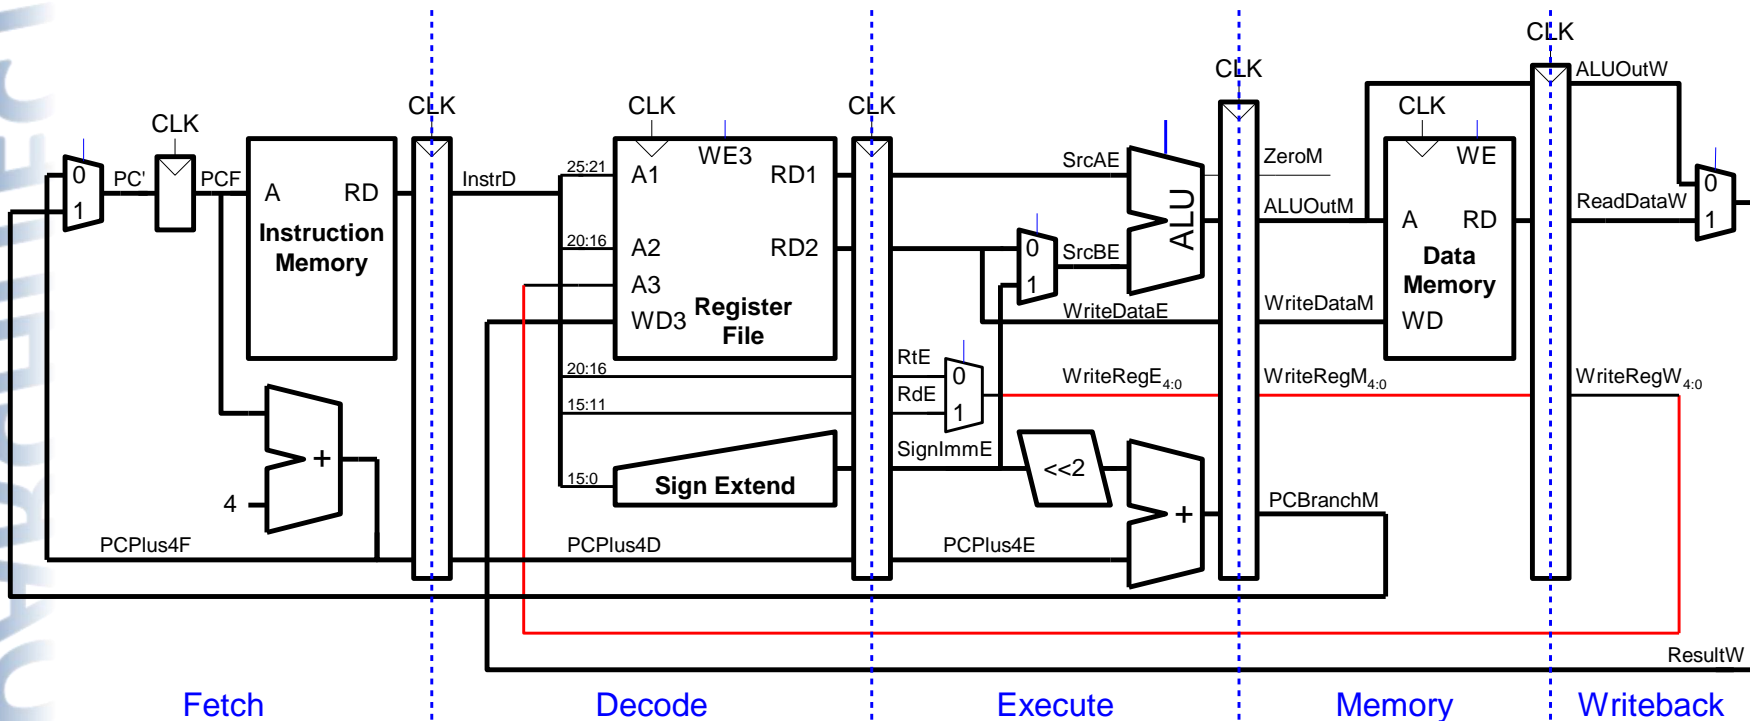
# Pipelined Processor Abstraction



# Single-Cycle & Pipelined Datapath

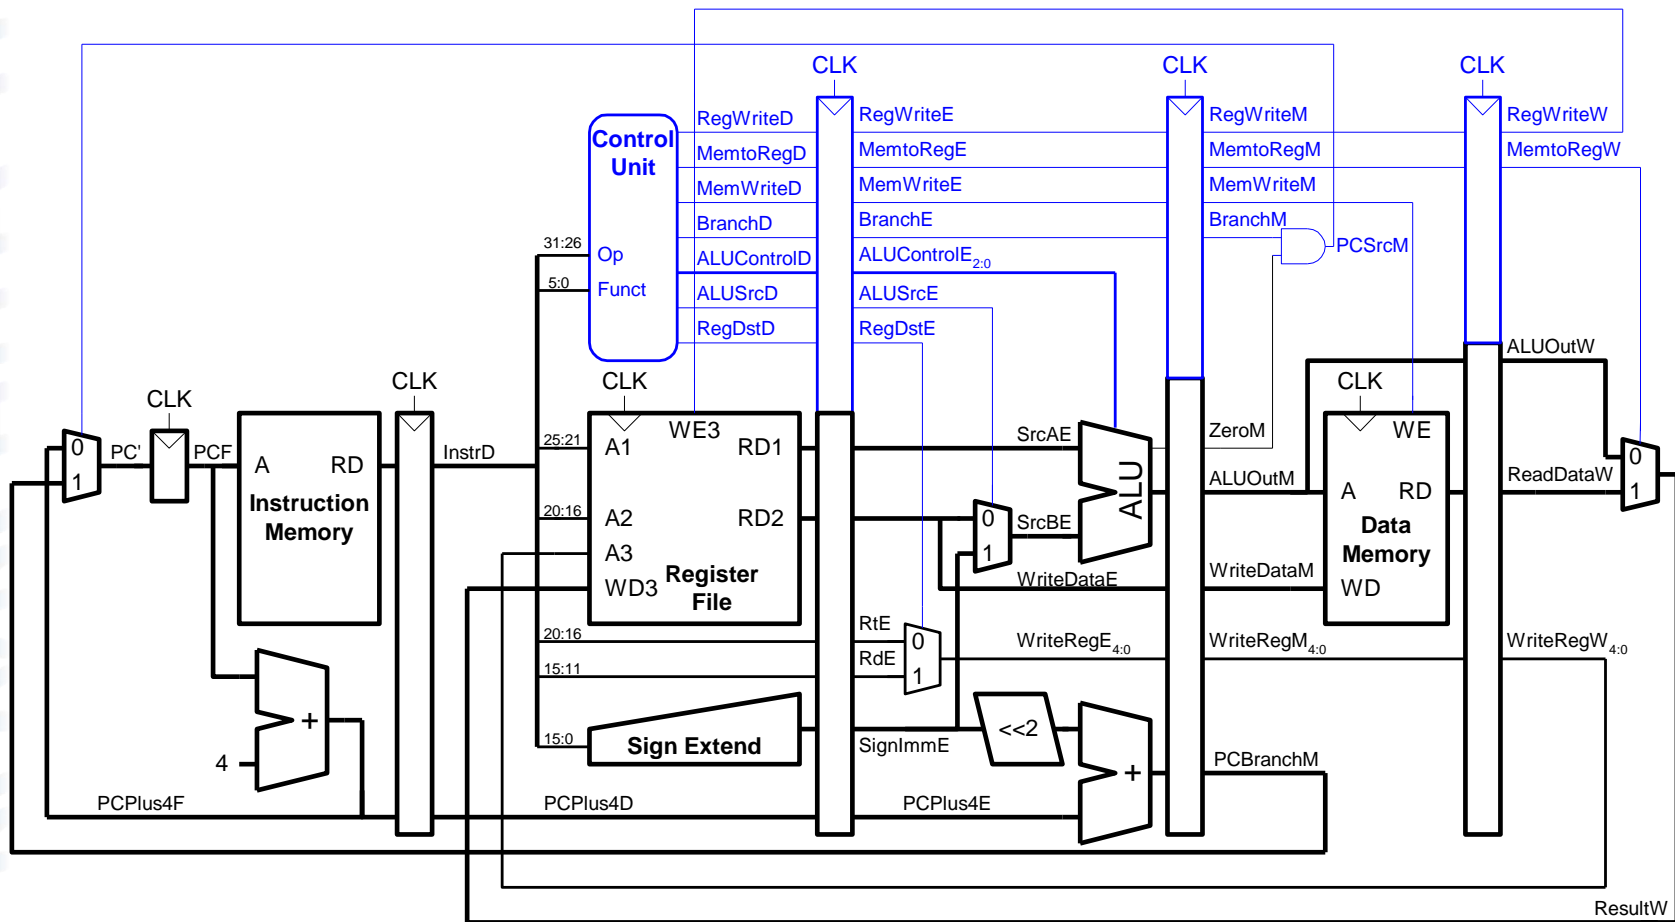


# Corrected Pipelined Datapath



*WriteReg* must arrive at same time as *Result*

# Pipelined Processor Control

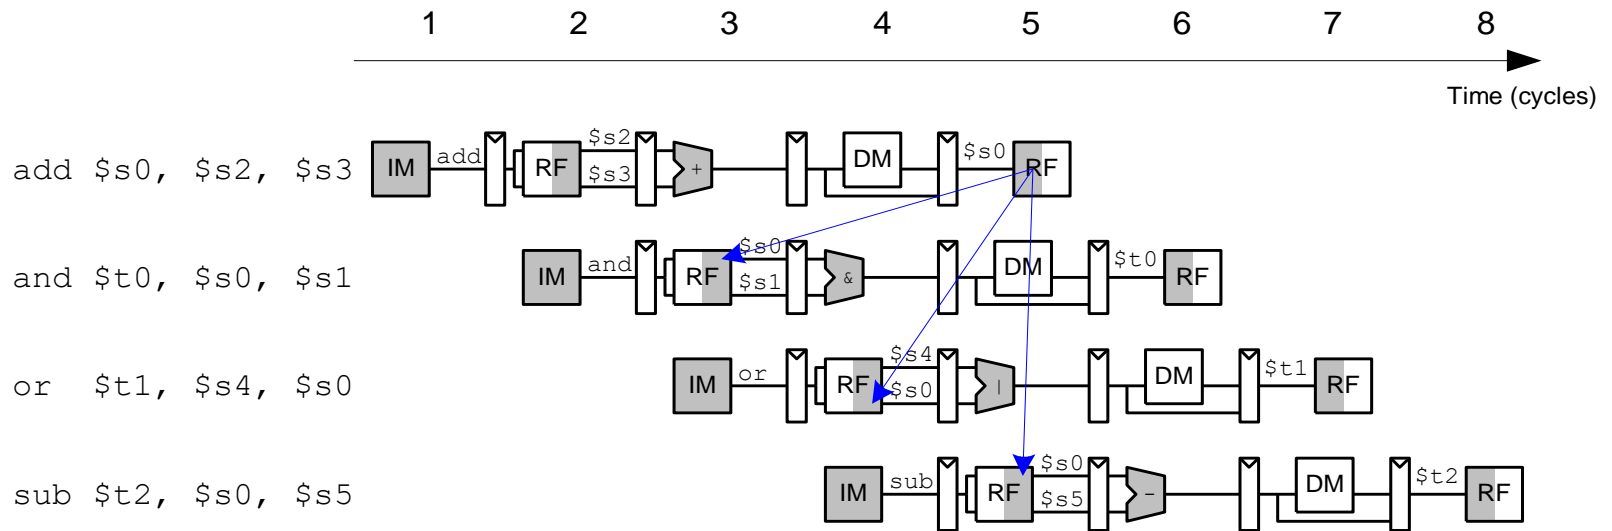


- Same control unit as single-cycle processor
- Control delayed to proper pipeline stage

# Pipeline Hazards

- When an instruction depends on result from instruction that hasn't completed
- Types:
  - **Data hazard:** register value not yet written back to register file
  - **Control hazard:** next instruction not decided yet (caused by branches) or target address not calculated yet (jumps and branches)

# Data Hazard



# Handling Data Hazards

## 2 SW fixes

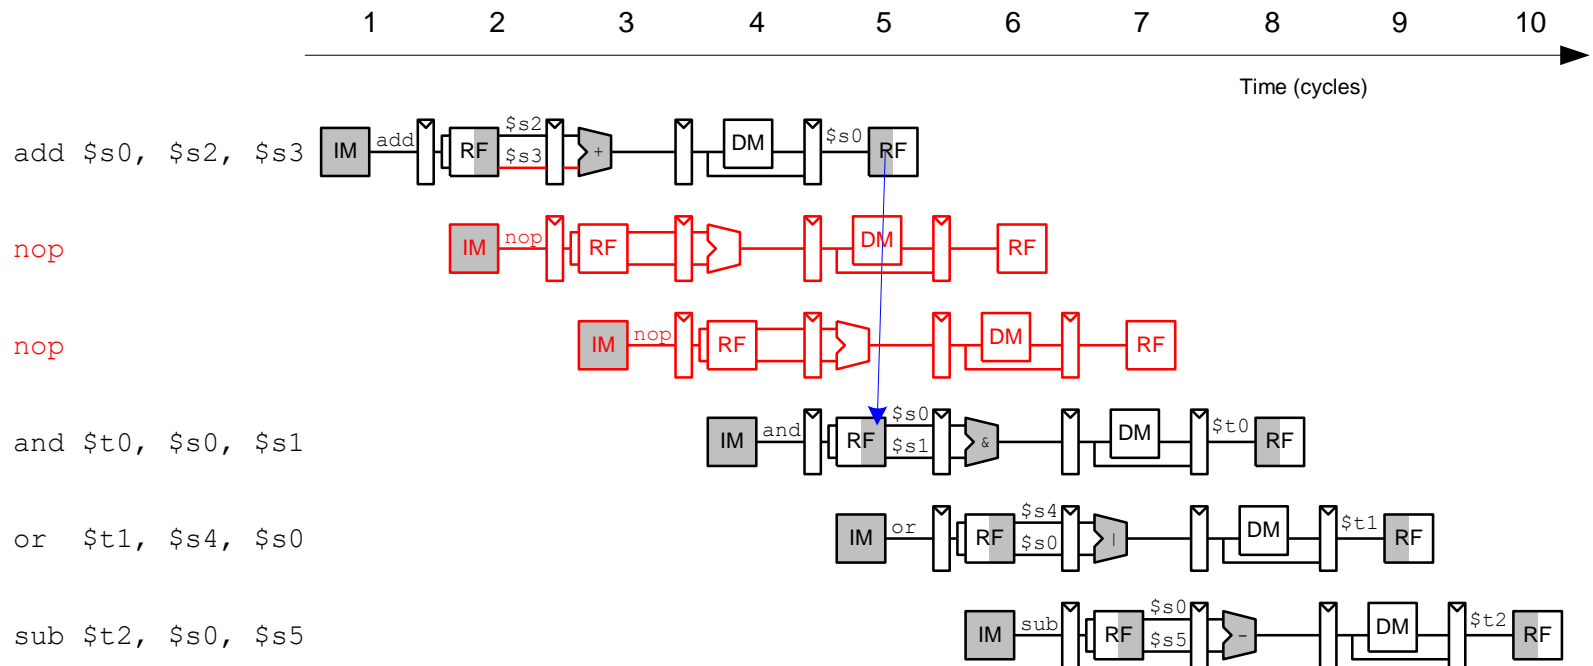
- Insert `nops` in code at compile time
- Rearrange code at compile time

## 2 HW fixes

- Stall the processor at run time
- Forward data at run time

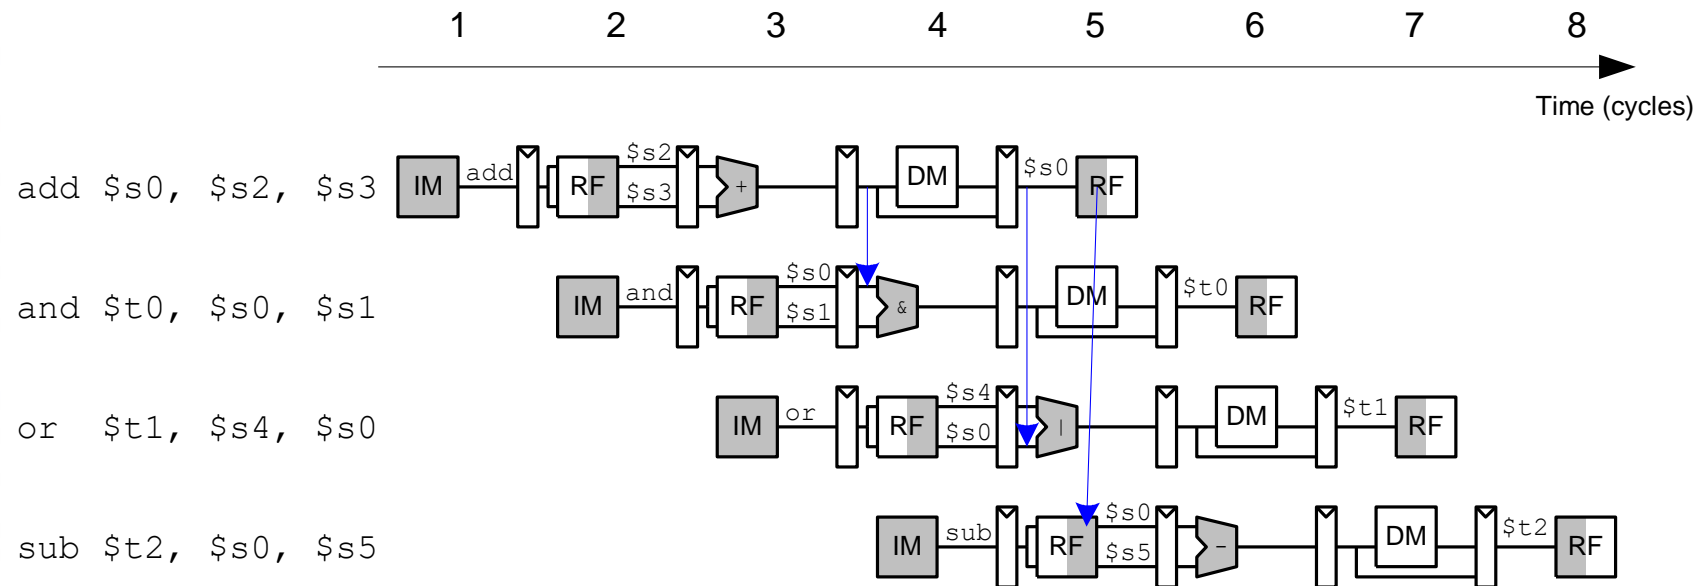
# Compile-Time Hazard Elimination

- Insert enough nops for result to be ready
- Or move independent useful instructions forward

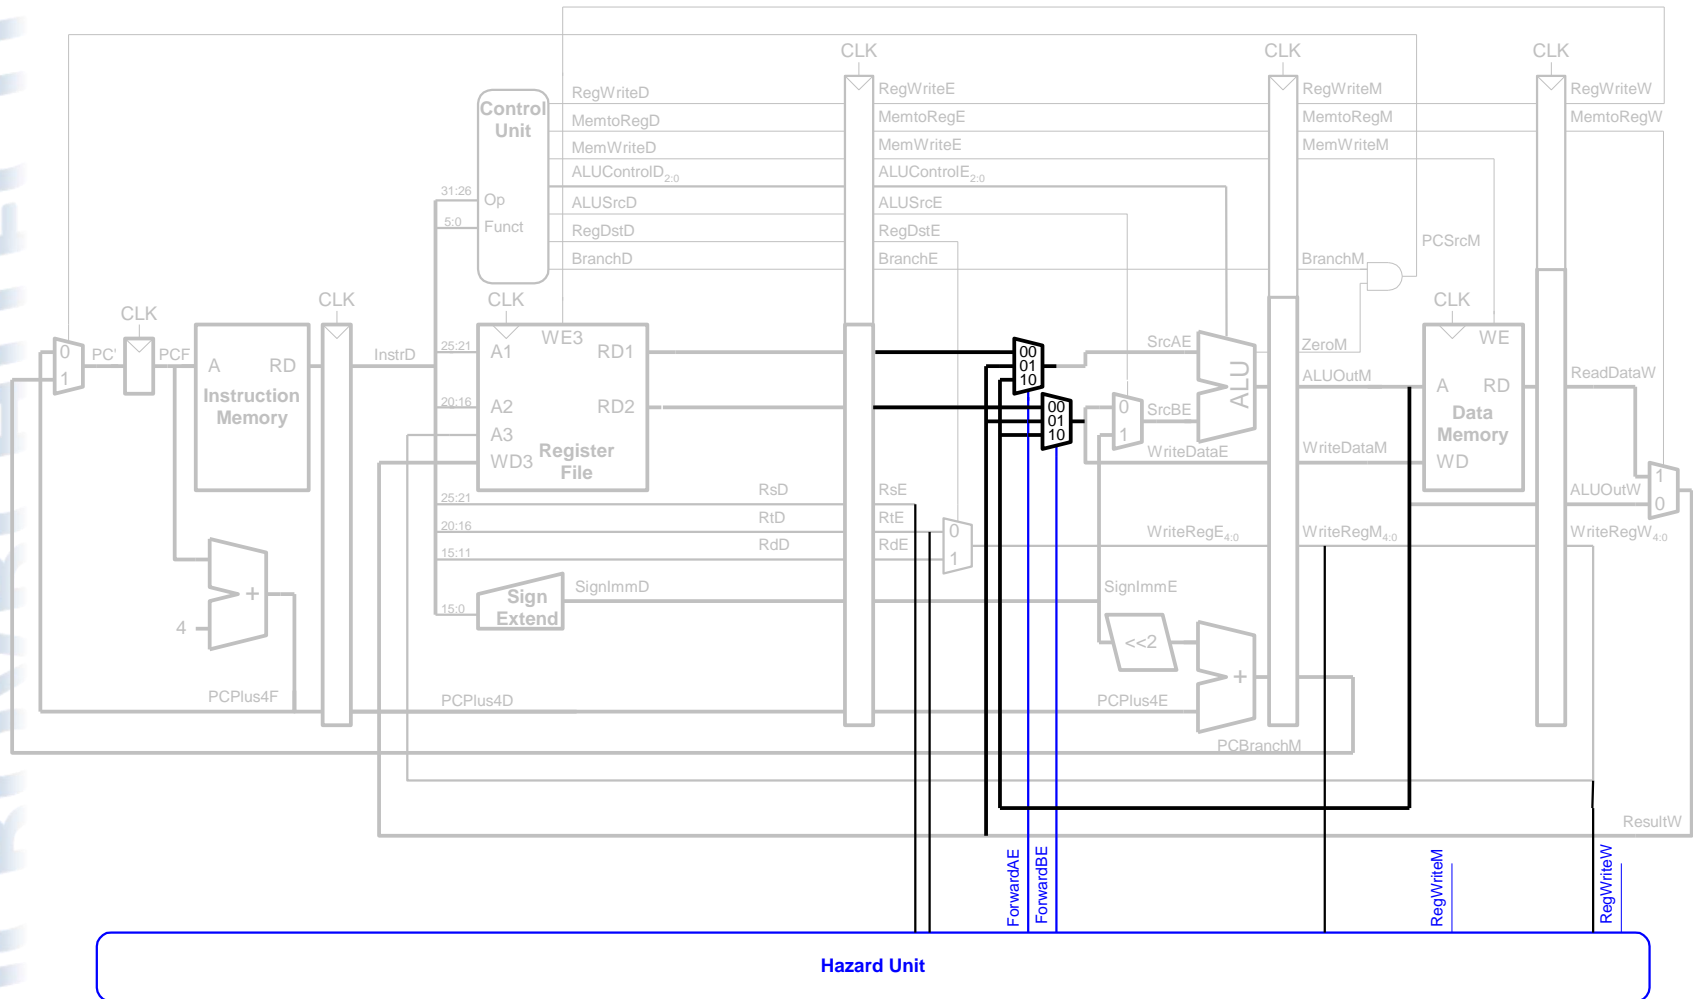




# Data Forwarding



# Data Forwarding



# Data Forwarding

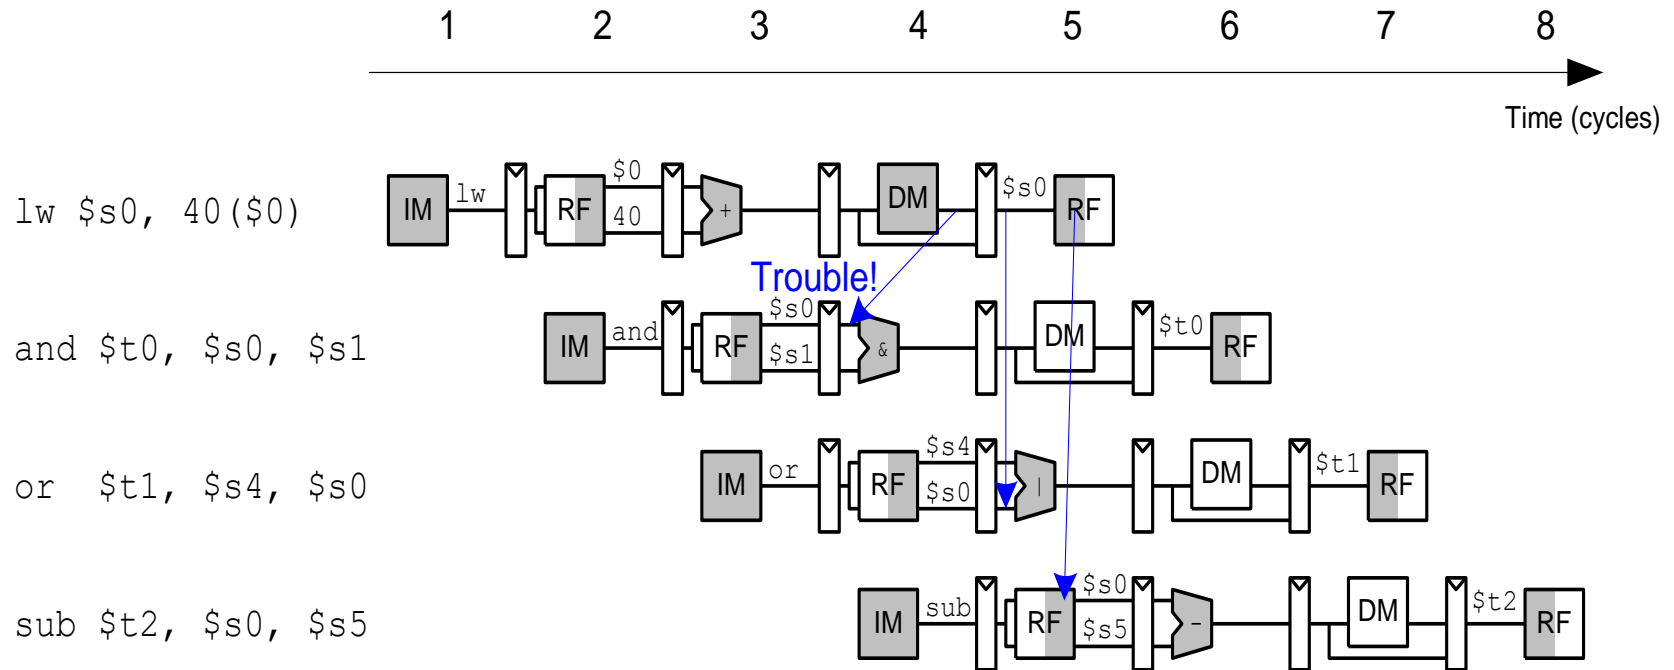
- Forward to Execute stage from either:
  - Memory stage or
  - Writeback stage
- Forwarding logic for *ForwardAE*:

```
if ((rsE != 0) AND (rsE == WriteRegM) AND RegWriteM)
    then ForwardAE = 10
else
    if ((rsE != 0) AND (rsE == WriteRegW) AND RegWriteW)
        then ForwardAE = 01
    else ForwardAE = 00
```

Forwarding logic for *ForwardBE* same, but replace *rsE* with *rtE*

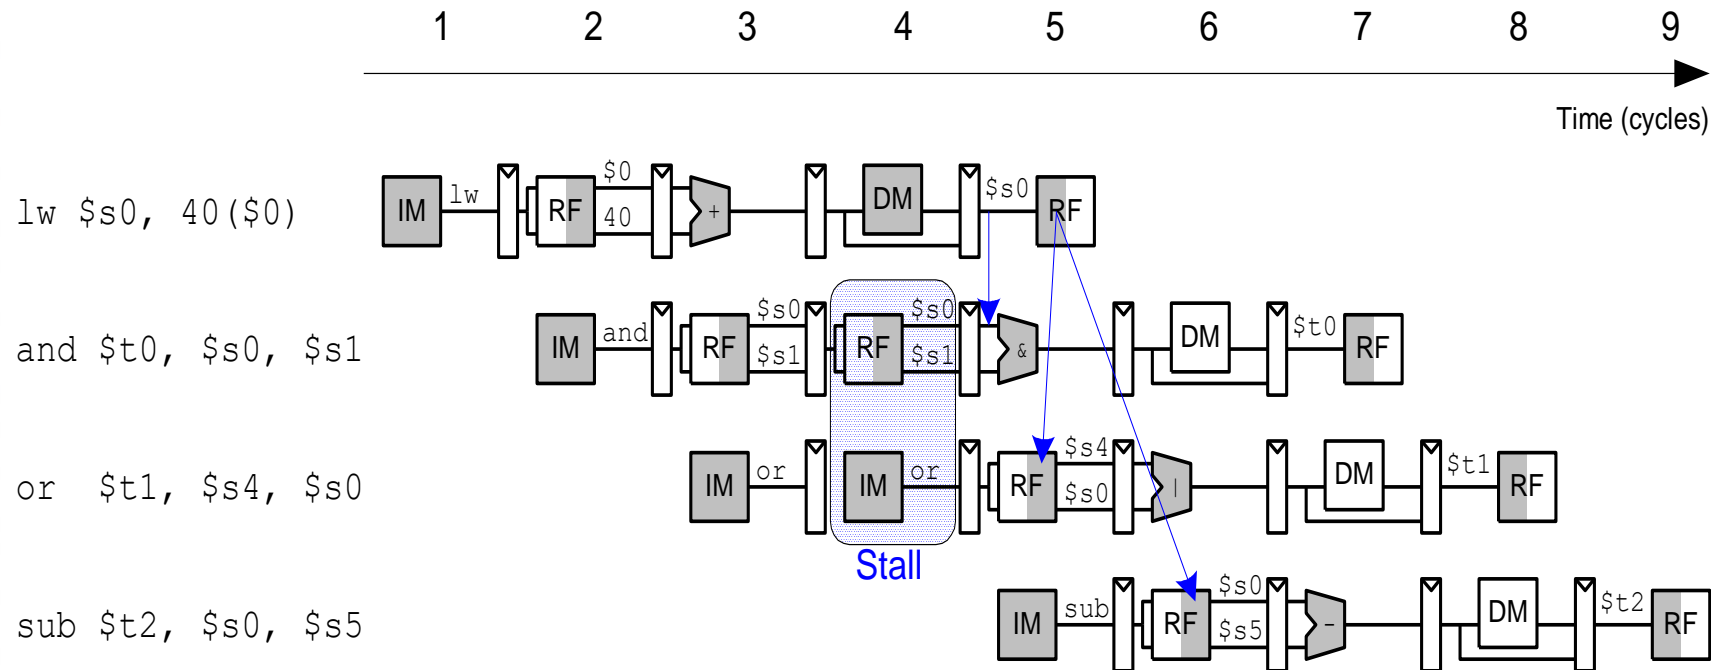
# Stalling

Forwarding on a load-use hazard isn't possible!




# Stalling

The HW solution is to stall the pipeline



# Stalling Hardware



# Stalling Logic

*lwstall* =

$((rsD == rtE) \text{ OR } (rtD == rtE)) \text{ AND } MemtoRegE$

*StallF* = *StallD* = *FlushE* = *lwstall*

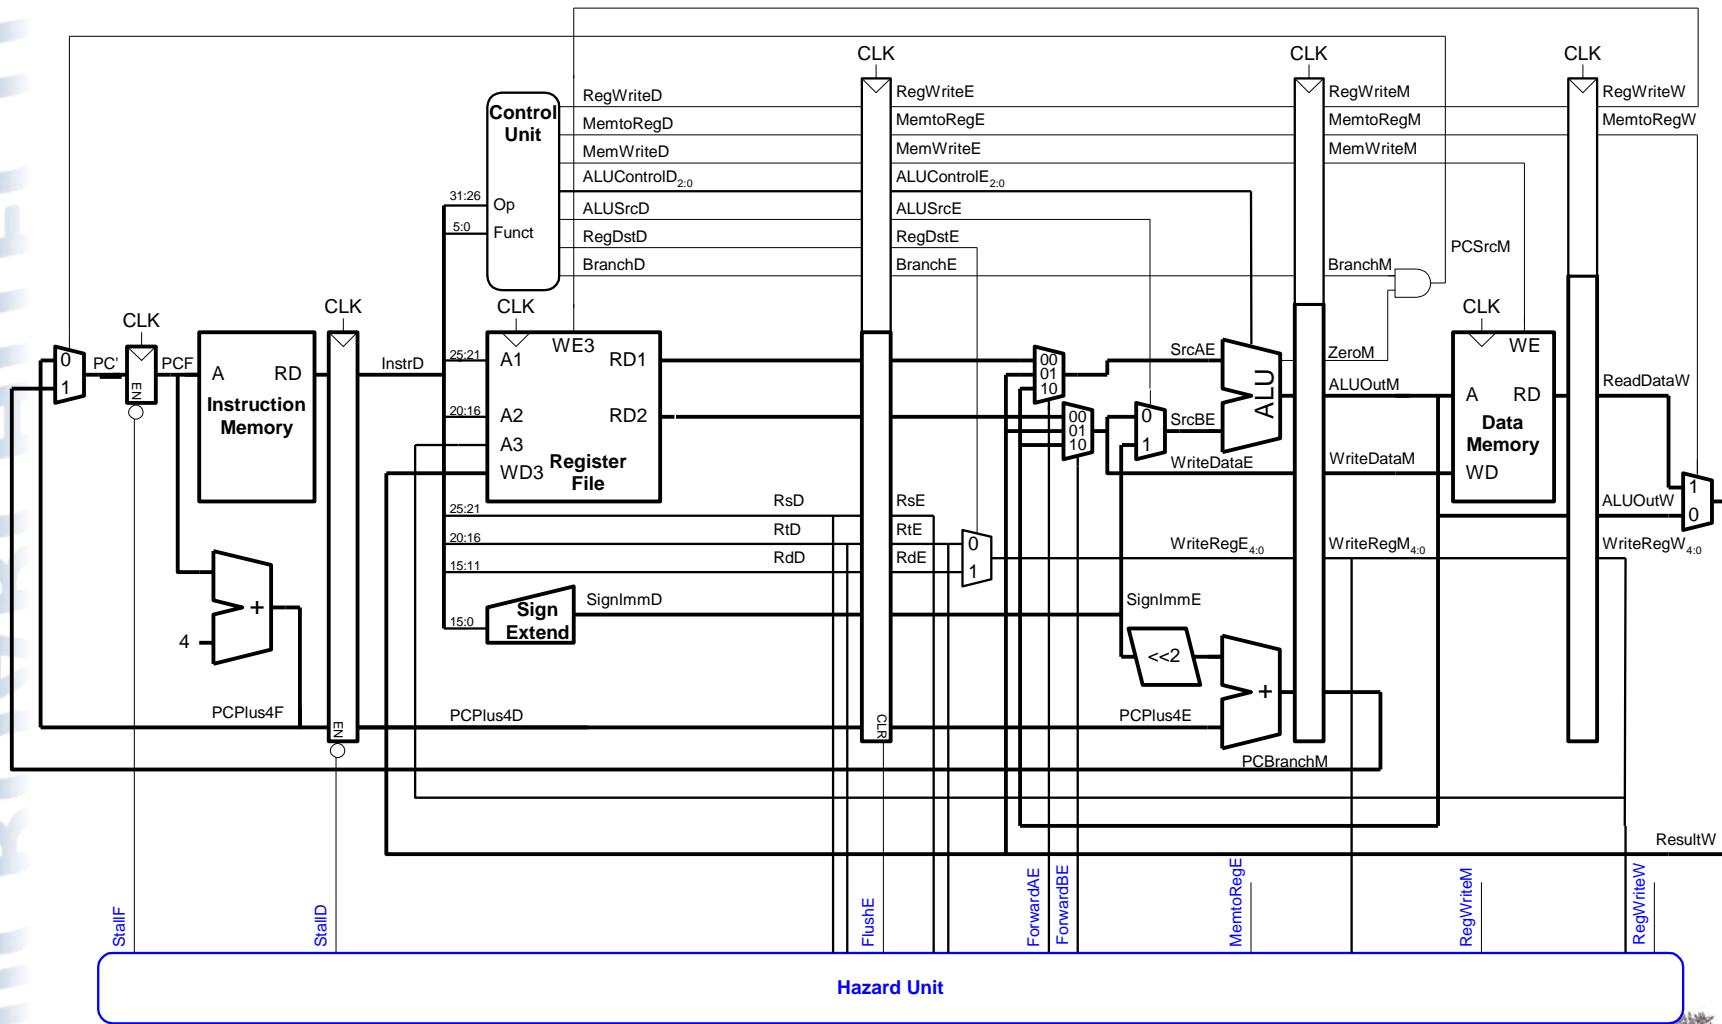
- By flushing the Execute stage, and stalling Fetch and Decode stages, the instruction flushed will simply be repeated in then next clock cycle, but this time with correct (forwarded) data!

# Control Hazards

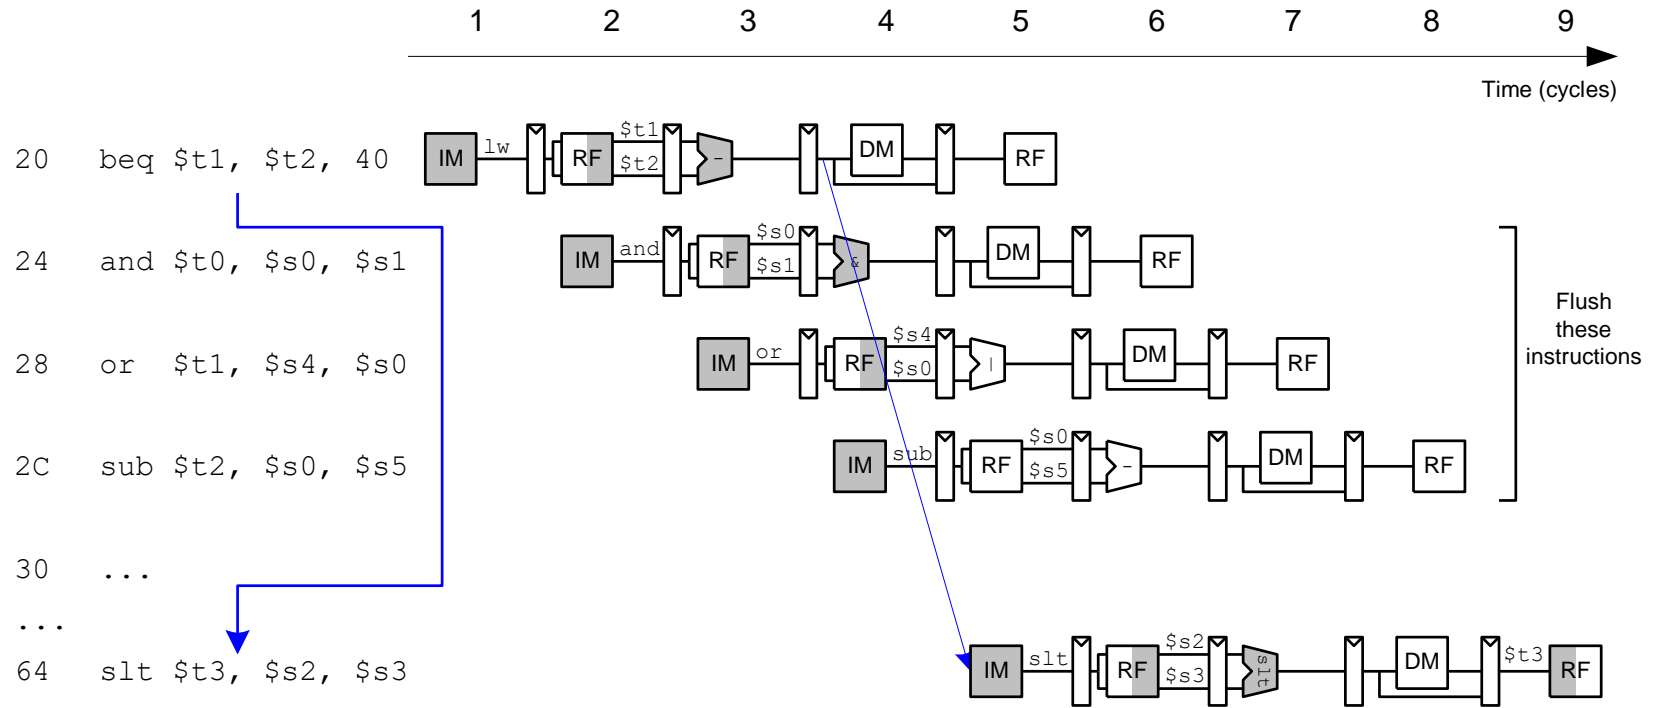
- **beq:**
  - branch not determined until 4<sup>th</sup> stage of pipeline
  - Instructions after the branch are fetched *before* the branch occurs
  - These instructions must be flushed if branch happens
- **Branch misprediction penalty**
  - the # of instruction flushed, when branch is taken
  - may be reduced by determining branch earlier




# Control Hazards: Original Pipeline



# Control Hazards

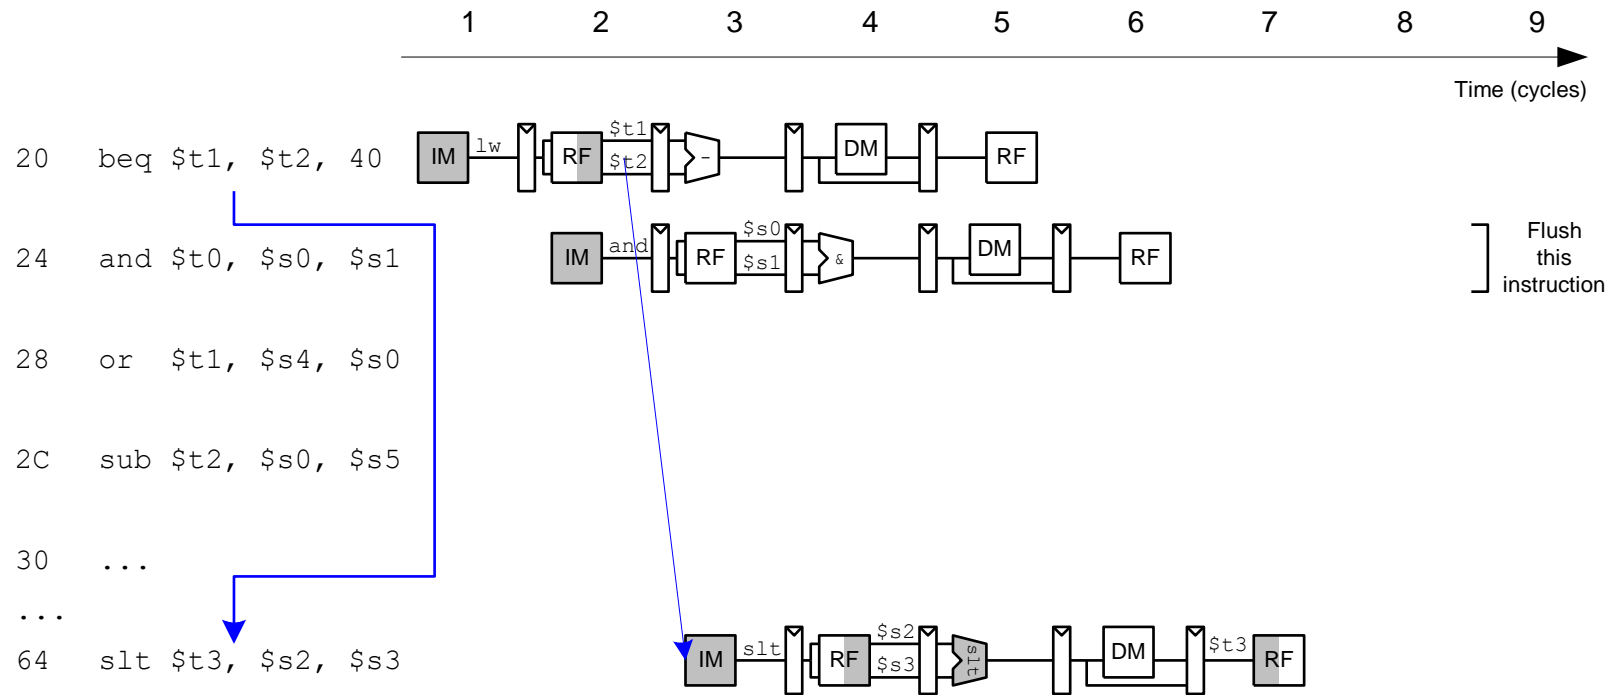


# Early Branch Resolution

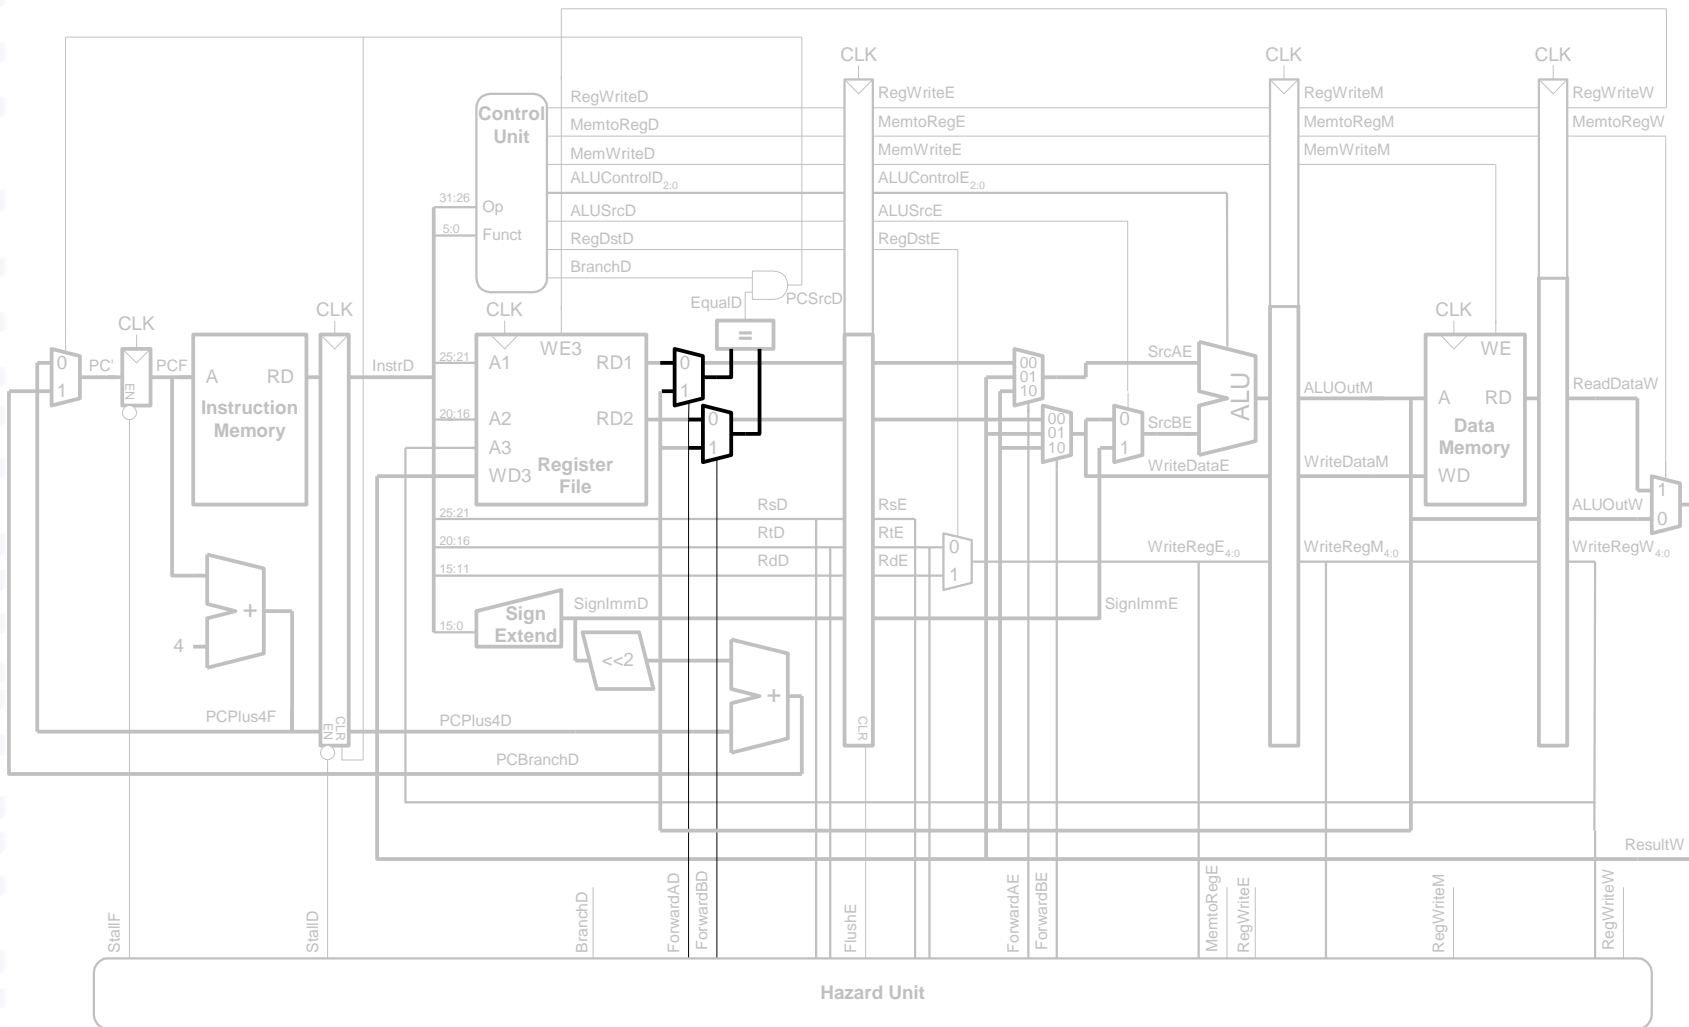


**But: introduced another data hazard in Decode stage!**

# Early Branch Resolution



# Forwarding to Early-branch HW



# Control Forwarding & Stalling Logic

- **Forwarding logic:**

**ForwardAD** = (*rsD* !=0) AND (*rsD* == *WriteRegM*) AND *RegWriteM*

**ForwardBD** = (*rtD* !=0) AND (*rtD* == *WriteRegM*) AND *RegWriteM*

- **Stalling logic:**

**branchstall** = *BranchD* AND *RegWriteE* AND  
(*WriteRegE* == *rsD* OR *WriteRegE* == *rtD*)

OR

*BranchD* AND *MemtoRegM* AND  
(*WriteRegM* == *rsD* OR *WriteRegM* == *rtD*)

**StallF** = **StallD** = **FlushE** = (*lwstall* OR *branchstall*)



# Branch Prediction

- Guess whether branch will be taken
  - Backward branches are usually taken (in bottom-tested loops)
  - Consider history to improve guess
- Good prediction significantly reduces fraction of branches requiring a flush
- Requires HW for history table, etc

# Pipelined Performance Example

- SPECINT2000 benchmark:
  - 25% loads
  - 10% stores
  - 11% branches
  - 2% jumps
  - 52% R-type
- Suppose:
  - 40% of loads used by next instruction
  - 25% of branches mispredicted
  - All jumps flush next instruction (JTA not ready)
- **What is the average CPI?**



# Calculation of Average CPI

- **Average CPI is the weighted average of  $CPI_{lw}$ ,  $CPI_{sw}$ ,  $CPI_{beq}$ ,  $CPI_j$  and  $CPI_{R-type}$**
- **For pipeline processors,  $CPI = 1 + \# \text{ of stall cycles}$**

Load CPI = 1 when no stall, = 2 when load-use occurs (1 stall)

–  $CPI_{lw} = 1(0.6) + 2(0.4) = 1.4$

–  $CPI_{sw} = 1$

Branch CPI = 1 when no stall, = 2 when it mispredicts and stalls

–  $CPI_{beq} = 1(0.75) + 2(0.25) = 1.25$

Jump CPI = 2 since it always requires 1 stall

–  $CPI_j = 2$

–  $CPI_{R-type} = 1$

$$\begin{aligned} \text{Average CPI} &= (0.25)(1.4) + (0.1)(1) + (0.11)(1.25) + (0.02)(2) + \\ &\quad (0.52)(1) = \mathbf{1.15} \end{aligned}$$



# Pipelined Performance

- Pipelined processor critical path:

$$T_c = \max \{$$
$$t_{pcq} + t_{mem} + t_{setup}$$
$$2(t_{RFread} + t_{mux} + t_{eq} + t_{AND} + t_{mux} + t_{setup})$$
$$t_{pcq} + t_{mux} + t_{mux} + t_{ALU} + t_{setup}$$
$$t_{pcq} + t_{memwrite} + t_{setup}$$
$$2(t_{pcq} + t_{mux} + t_{RFwrite}) \}$$

# Pipelined Performance Example

Element	Parameter	Delay (ps)
Register clock-to-Q	$t_{pcq\_PC}$	30
Register setup	$t_{setup}$	20
Multiplexer	$t_{mux}$	25
ALU	$t_{ALU}$	200
Memory read	$t_{mem}$	250
Register file read	$t_{RFread}$	150
Register file setup	$t_{RFsetup}$	20
Equality comparator	$t_{eq}$	40
AND gate	$t_{AND}$	15
Memory write	$T_{memwrite}$	220
Register file write	$t_{RFwrite}$	100 ps

$$\begin{aligned}
 T_c &= 2(t_{RFread} + t_{mux} + t_{eq} + t_{AND} + t_{mux} + t_{setup}) \\
 &= 2[150 + 25 + 40 + 15 + 25 + 20] \text{ ps} = \mathbf{550 \text{ ps}}
 \end{aligned}$$

# Pipelined Performance Example

Program with IC = 100 billion instructions

$$\begin{aligned}\text{Execution Time} &= \text{IC} \times \text{CPI} \times T_c \\ &= (100 \times 10^9)(1.15)(550 \times 10^{-12}) \\ &= \mathbf{63 \text{ seconds}}\end{aligned}$$

# Processor Performance Comparison

Processor	Execution Time (seconds)	Speedup (single-cycle as baseline)
<b>Single-cycle</b>	92.5	1
<b>Multicycle</b>	133	0.70
<b>Pipelined</b>	63	1.47

# Advanced Microarchitecture

- Deep Pipelining
- Branch Prediction
- Superscalar Processors
- Out of Order Processors
- Register Renaming
- SIMD
- Multithreading
- Multiprocessors

# Deep Pipelining

- 10-20 stages typical
- Number of stages limited by:
  - Pipeline hazards
  - Sequencing overhead
  - Power
  - Cost

# Branch Prediction

- Ideal pipelined processor:  $CPI = 1$
- Branch misprediction increases CPI
- **Static branch prediction:**
  - Check direction of branch (forward or backward)
  - If backward, predict taken
  - Else, predict not taken
- **Dynamic branch prediction:**
  - Keep history of last (several hundred) branches in *branch target buffer*, record:
    - Branch destination
    - Whether branch was taken



# Branch Prediction Example


```
add    $s1, $0, $0          # sum = 0
add    $s0, $0, $0          # i   = 0
addi   $t0, $0, 10          # $t0 = 10

for:
    beq  $s0, $t0, done      # if i == 10, branch
    add  $s1, $s1, $s0       # sum = sum + i
    addi $s0, $s0, 1         # increment i
    j    for
done:
```

# 1-Bit Branch Predictor

- Remembers whether branch was taken the last time and does the same thing
- Mispredicts first and last branch of loop

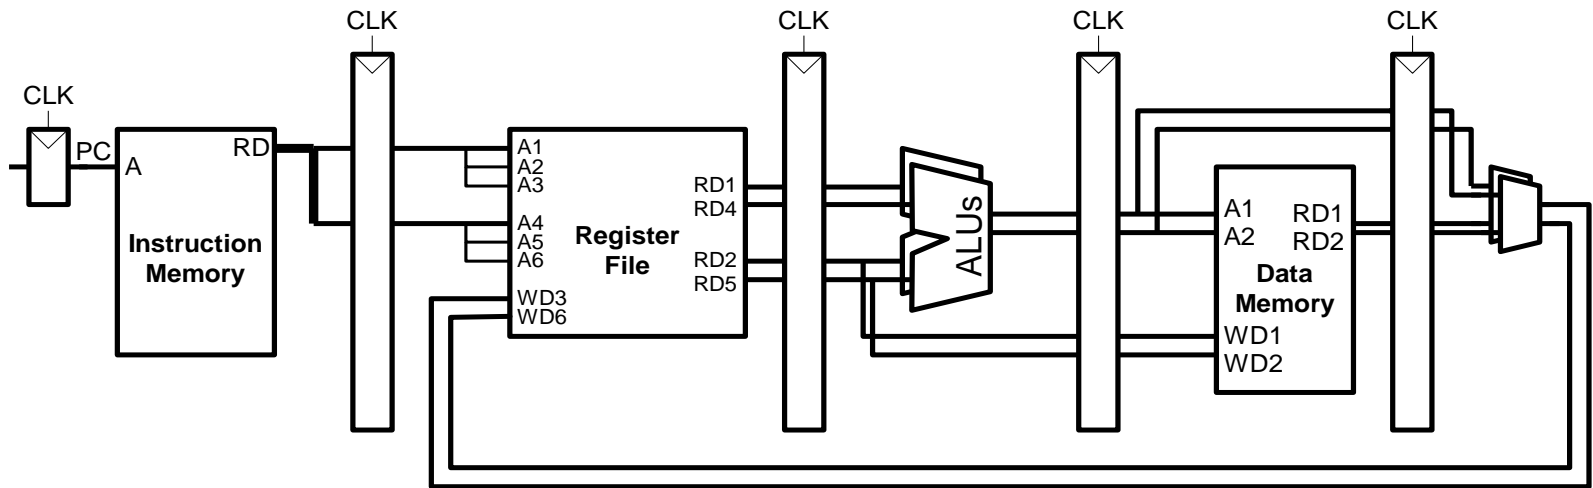
# 2-Bit Branch Predictor



**Only mispredicts the last branch of the loop**

# Superscalar

- Multiple copies of datapath execute multiple instructions at once
- Dependencies make it tricky to issue multiple instructions at once

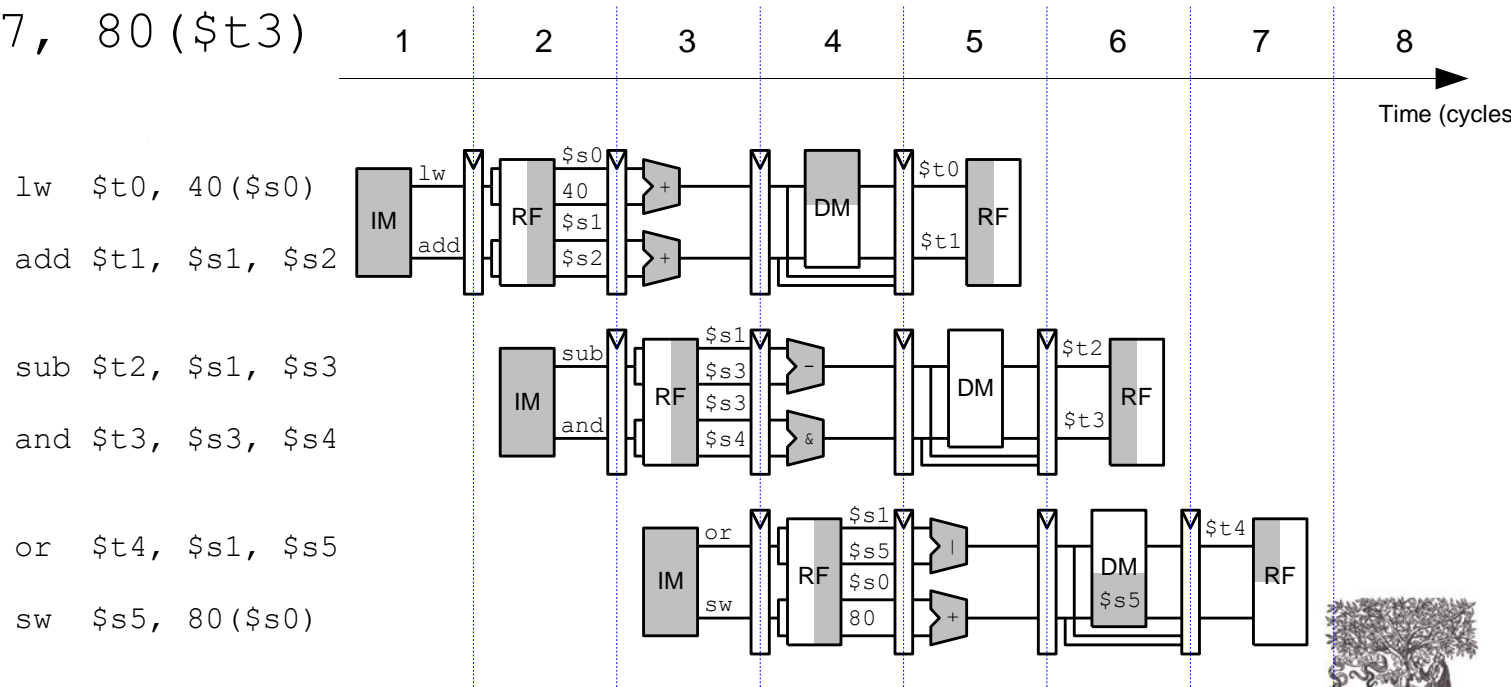


# Superscalar Example

```
lw    $t0, 40($s0)
add   $t1, $t0, $s1
sub   $t0, $s2, $s3
and   $t2, $s4, $t0
or    $t3, $s5, $s6
sw    $s7, 80($t3)
```

**Ideal IPC: 2**

**Actual IPC: 2**



# Superscalar with Dependencies

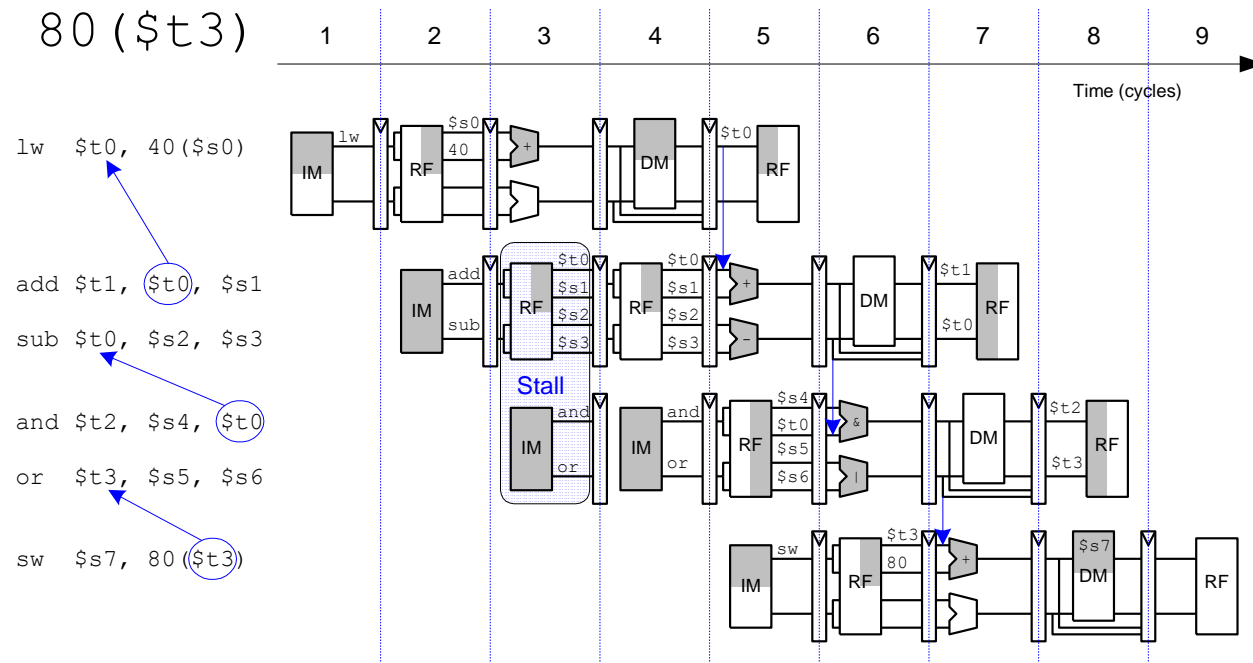
```
lw    $t0, 40($s0)
add   $t1, $t0, $s1
sub   $t0, $s2, $s3
and   $t2, $s4, $t0
or    $t3, $s5, $s6
sw    $s7, 80($t3)
```

**Ideal IPC:**

**2**

**Actual IPC:**

**6/5 = 1.17**



# Out of Order Processor

- Looks ahead across multiple instructions
- Issues as many instructions as possible at once
- Issues instructions out of order (as long as no dependencies)
- **Dependencies:**
  - **RAW** (read after write): one instruction writes, later instruction reads a register
  - **WAR** (write after read): one instruction reads, later instruction writes a register
  - **WAW** (write after write): one instruction writes, later instruction writes a register

# Out of Order Processor

- **Instruction level parallelism (ILP):** number of instruction that can be issued simultaneously (average  $< 3$ )
- **Scoreboard:** table that keeps track of:
  - Instructions waiting to issue
  - Available functional units
  - Dependencies



# Out of Order Processor Example

```
lw    $t0, 40($s0)
add   $t1, $t0, $s1
sub   $t0, $s2, $s3
and   $t2, $s4, $t0
or    $t3, $s5, $s6
sw    $s7, 80($t3)
```

Ideal IPC:

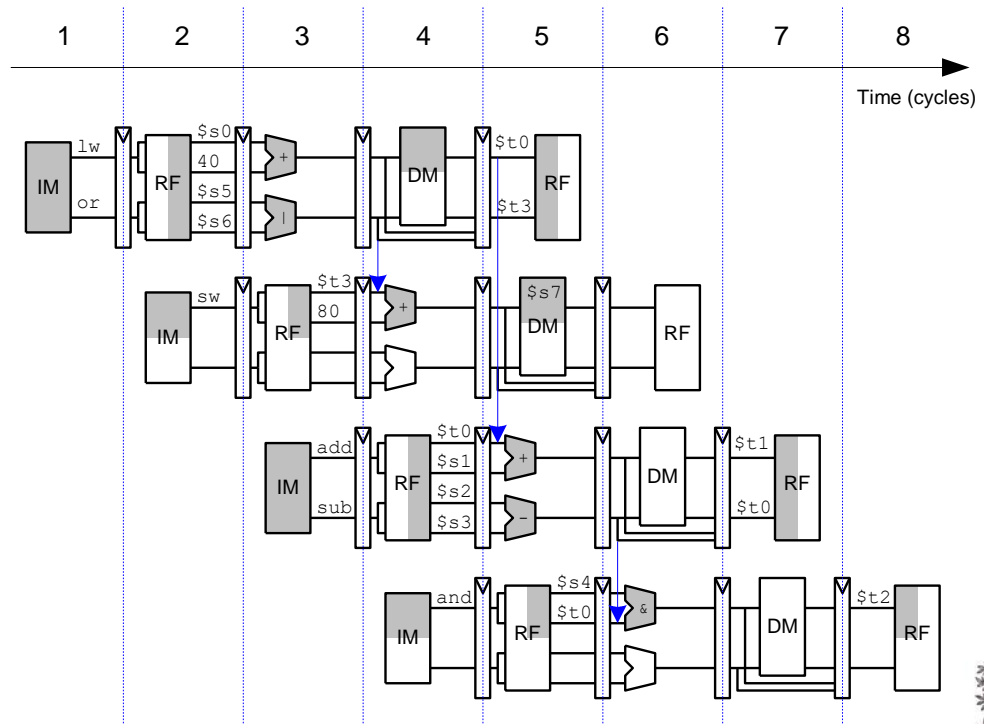
2

Actual IPC:

$6/4 = 1.5$

lw \$t0, 40(\$s0)  
 or \$t3, \$s5, \$s6  
 sw \$s7, 80(\$t3)  
 add \$t1, \$t0, \$s1  
 sub \$t0, \$s2, \$s3  
 and \$t2, \$s4, \$t0

RAW  
 two cycle latency  
 between load and  
 use of \$t0  
 RAW  
 WAR  
 RAW



# Register Renaming

```
lw    $t0, 40($s0)
add   $t1, $t0, $s1
sub   $t0, $s2, $s3
and   $t2, $s4, $t0
or    $t3, $s5, $s6
sw    $s7, 80($t3)
```

Ideal IPC:

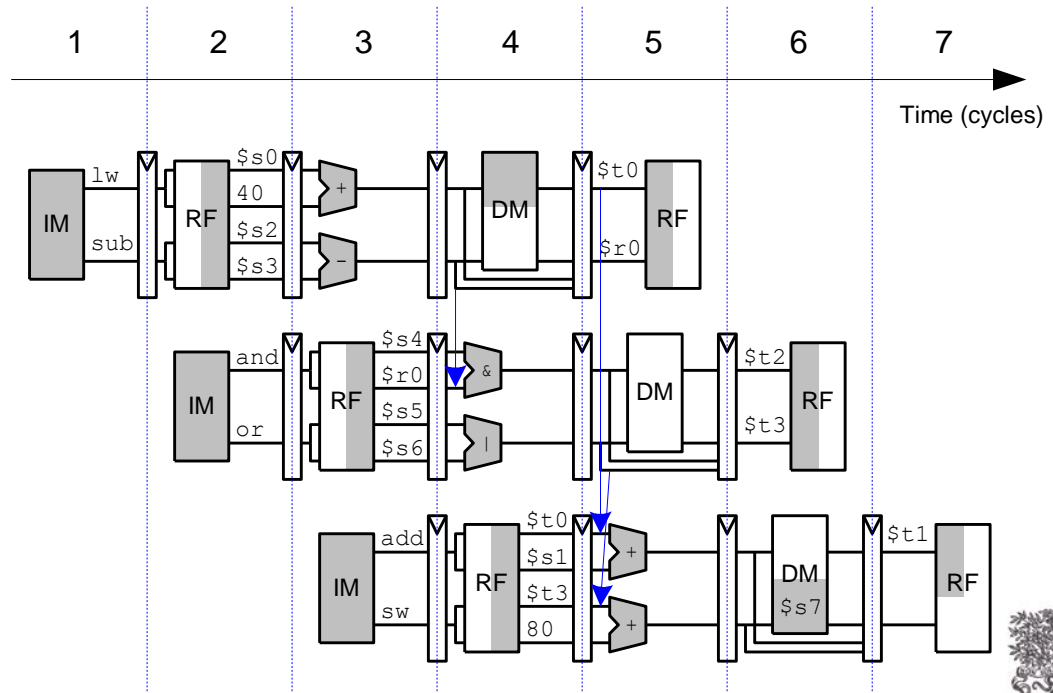
2

Actual IPC:

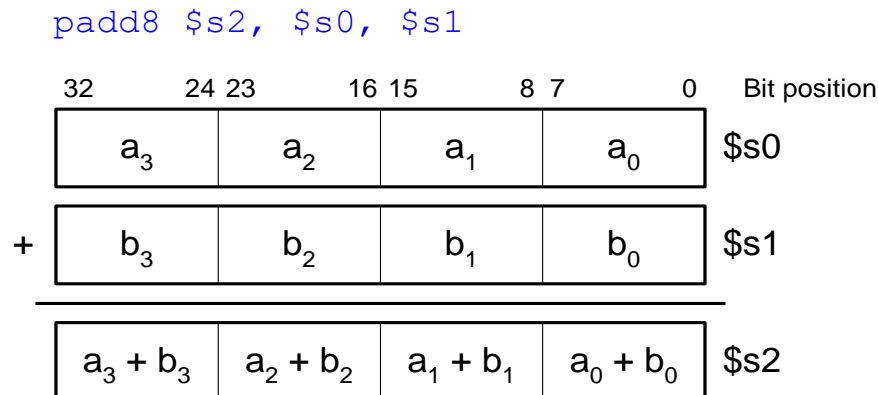
$6/3 = 2$

lw \$t0, 40(\$s0)  
 sub \$r0, \$s2, \$s3  
 and \$t2, \$s4, \$r0  
 or \$t3, \$s5, \$s6  
 add \$t1, \$t0, \$s1  
 sw \$s7, 80(\$t3)

2-cycle RAW  
 RAW  
 RAW



- Single Instruction Multiple Data (SIMD)
  - Single instruction acts on multiple pieces of data at once
  - Common application: graphics
  - Perform short arithmetic operations (also called *packed arithmetic*)
- For example, add four 8-bit elements



# Advanced Architecture Techniques

- **Multithreading**
  - Wordprocessor: thread for typing, spell checking, printing
- **Multiprocessors**
  - Multiple processors (cores) on a single chip

# Threading: Definitions

- **Process:** program running on a computer
  - Multiple processes can run at once: e.g., surfing Web, playing music, writing a paper
- **Thread:** part of a program
  - Each process has multiple threads: e.g., a word processor may have threads for typing, spell checking, printing



# Threads in Conventional Processor

- One thread runs at once
- When one thread stalls (for example, waiting for memory):
  - Architectural state of that thread stored
  - Architectural state of waiting thread loaded into processor and it runs
  - Called **context switching**
- Appears to user like all threads running simultaneously

# Multithreading

- Multiple copies of architectural state
- Multiple threads **active** at once:
  - When one thread stalls, another runs immediately
  - If one thread can't keep all execution units busy, another thread can use them
- Does not increase instruction-level parallelism (ILP) of single thread, but increases throughput

Intel calls this “hyperthreading”

# Multiprocessors

- Multiple processors (cores) with a method of communication between them
- Types:
  - **Homogeneous:** multiple cores with shared memory
  - **Heterogeneous:** separate cores for different tasks (for example, DSP and CPU in cell phone)
  - **Clusters:** each core has own memory system